

```
In [0]:
```

```
from google.colab import drive  
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call  
drive.mount("/content/drive", force_remount=True).
```

```
In [0]:
```

```
%cd drive/My Drive
```

```
[Errno 2] No such file or directory: 'drive/My Drive'  
/content/drive/My Drive/Assignments_DonorsChoose_2018
```

```
In [0]:
```

```
%cd Assignments_DonorsChoose_2018
```

```
[Errno 2] No such file or directory: 'Assignments_DonorsChoose_2018'  
/content/drive/My Drive/Assignments_DonorsChoose_2018
```

```
In [0]:
```

```
%matplotlib inline  
import warnings  
warnings.filterwarnings("ignore")  
  
import sqlite3  
import pandas as pd  
import numpy as np  
import nltk  
import string  
import matplotlib.pyplot as plt  
import seaborn as sns  
import sklearn  
from sklearn.feature_extraction.text import TfidfTransformer  
from sklearn.feature_extraction.text import TfidfVectorizer  
  
from sklearn.feature_extraction.text import CountVectorizer  
from sklearn.metrics import confusion_matrix  
from sklearn import metrics  
from sklearn.metrics import roc_curve, auc  
from nltk.stem.porter import PorterStemmer  
  
import re  
# Tutorial about Python regular expressions: https://pymotw.com/2/re/  
import string  
from nltk.corpus import stopwords  
from nltk.stem import PorterStemmer  
from nltk.stem.wordnet import WordNetLemmatizer  
  
from gensim.models import Word2Vec  
from gensim.models import KeyedVectors  
import pickle  
  
from tqdm import tqdm  
import os  
  
from plotly import plotly  
import plotly.offline as offline  
import plotly.graph_objs as go  
offline.init_notebook_mode()  
from collections import Counter
```

```
In [0]:
```

```
project_data = pd.read_csv('train_data.csv')
```

```
resource_data = pd.read_csv('resources.csv')
```

In [0]:

```
# how to replace elements in list python: https://stackoverflow.com/a/2582163/4084039
cols = ['Date' if x=='project_submitted_datetime' else x for x in list(project_data.columns)]

#sort dataframe based on time pandas python: https://stackoverflow.com/a/49702492/4084039
project_data['Date'] = pd.to_datetime(project_data['project_submitted_datetime'])
project_data.drop('project_submitted_datetime', axis=1, inplace=True)
project_data.sort_values(by=['Date'], inplace=True)

# how to reorder columns pandas python: https://stackoverflow.com/a/13148611/4084039
project_data = project_data[cols]
#project_data.head(2)
```

In [0]:

```
price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset_index()
#price_data.head(2)
```

In [0]:

```
project_data = pd.merge(project_data, price_data, on='id', how='left')
```

1.2 preprocessing of project_subject_categories

In [0]:

```
categories = list(project_data['project_subject_categories'].values)
# remove special characters from list of strings python:
https://stackoverflow.com/a/47301924/4084039

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python
cat_list = []
for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','):
        # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split():
            # this will split each of the category based on space "Math & Science"=> "Math", "&", "Science"
            j=j.replace('The', '') # if we have the words "The" we are going to replace it with ''(i.e removing 'The')
        j = j.replace(' ', '') # we are placeing all the ' ' (space) with ''(empty) ex:"Math & Science"=>"Math&Science"
        temp+=j.strip()+" "
    temp = temp.replace('&', '_') # we are replacing the & value into
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))
```

1.3 preprocessing of project_subject_subcategories

In [0]:

```
sub_categories = list(project_data['project_subject_subcategories'].values)
# remove special characters from list of strings python.
```

```

# remove special characters from list of strings python:
https://stackoverflow.com/a/47301924/4084039

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

sub_cat_list = []
for i in sub_categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','):# it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"=> "Math", "&", "Science"
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ''(i.e removing 'The')
            j = j.replace(' ', '') # we are placing all the ''(space) with ''(empty) ex:"Math & Science"=>"Math&Science"
            temp +=j.strip()+" "# abc ".strip() will return "abc", remove the trailing spaces
            temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/4084039
my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

1.3 Text preprocessing

In [0]:

```

# merge two column text dataframe:
project_data["essay"] = project_data["project_essay_1"].map(str) + \
                        project_data["project_essay_2"].map(str) + \
                        project_data["project_essay_3"].map(str) + \
                        project_data["project_essay_4"].map(str)

```

In [0]:

```

# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"\n\t", " not", phrase)
    phrase = re.sub(r"\re", " are", phrase)
    phrase = re.sub(r"\s", " is", phrase)
    phrase = re.sub(r"\d", " would", phrase)
    phrase = re.sub(r"\ll", " will", phrase)
    phrase = re.sub(r"\t", " not", phrase)
    phrase = re.sub(r"\ve", " have", phrase)
    phrase = re.sub(r"\m", " am", phrase)
    return phrase

```

In [0]:

```

# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've",
\           "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his',
'himself', \

```

```
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', \
    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", \
'these', 'those', \
    'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', \
'do', 'does', \
    'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', \
'while', 'of', \
    'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', \
'before', 'after', \
    'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', \
', again', 'further', \
    'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', \
'few', 'more', \
    'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', \
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', \
', m', 'o', 're', \
    've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn", "doesn't", \
'hadn', \
    'hadn't', 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', \
"mightn't", 'mustn', \
    'mustn't', 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', \
"wasn't", 'weren', "weren't", \
    'won', "won't", 'wouldn', "wouldn't"]
```

In [0]:

```
def find_num(text):
    if re.findall(r'\d+', text):
        return 1
    return 0

project_data['numerical_digits'] = project_data['project_resource_summary'].apply(lambda x: find_num(x))
```

In [0]:

```
project_data['project_grade_category']=project_data['project_grade_category'].str.replace(' ', '_')
project_data['project_grade_category']=project_data['project_grade_category'].str.replace('-', 'to')
set(project_data['project_grade_category'])
```

Out[0]:

```
{'Grades_3to5', 'Grades_6to8', 'Grades_9to12', 'Grades_PreKto2'}
```

In [0]:

```
project_data['teacher_prefix']=project_data['teacher_prefix'].fillna('Mrs.')
```

In [0]:

```
#project_data = project_data[:50000]
```

In [0]:

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import accuracy_score
from collections import Counter
import matplotlib.pyplot as plt
import numpy as np
from scipy.sparse import csr_matrix
import time

project_data_features = project_data.copy()
project_data_features.drop('project_is_approved', axis=1, inplace=True)
y=list(project_data['project_is_approved'])
X_train, X_test, y_train, y_test = train_test_split(project_data_features, y, stratify=y, test_size=0.25)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, stratify=y_train, test_size=0.2
5)
```

```
In [0]:
```

```
X_train.shape, X_val.shape, X_test.shape
```

```
Out[0]:
```

```
((61452, 20), (20484, 20), (27312, 20))
```

```
In [0]:
```

```
u=[[],[],[],[],[],[]]
def prob_encode(df):
    for i,col in enumerate(['teacher_prefix', 'school_state', 'clean_categories',
'clean_subcategories', 'project_grade_category','numerical_digits']):
        df[col+str(0)] = 0.5
        df[col+str(1)] = 0.05
        #print(len(u[i]),i)
    if len(u[i])==0:
        u[i]=df.groupby(col).project_is_approved.apply(lambda g: g.value_counts() /len(g) ).unstack().stack()
    ind=df.columns.get_loc(col)
    for row in df.itertuples():
        try:
            if u[i][row[ind+1]].any():
                df.at[row.Index, col+'0'] = list(u[i][row[ind+1]])[0]
                df.at[row.Index, col+'1'] = list(u[i][row[ind+1]])[1]
        except:
            df.at[row.Index, col+'0'] = 0.5
            df.at[row.Index, col+'1'] = 0.05
    i=i+1
```

```
In [0]:
```

```
X_train['project_is_approved'] = y_train
```

```
In [0]:
```

```
import time
start = time.time()
prob_encode(X_train)
prob_encode(X_test)
prob_encode(X_val)
X_train.drop(X_train.columns[len(X_train.columns)-1], axis=1, inplace=True)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```

```
execution time in minutes: 13.099024506409963
execution time in hours: 0
```

```
In [0]:
```

```
from sklearn.preprocessing import StandardScaler
def standardize_data(df_tr,df_cv,df_te,column_name):
    standardized_vec = StandardScaler(with_mean=False)
    # here it will learn mu and sigma
    standardized_vec.fit(df_tr[column_name].values.reshape(-1,1))

    # with the learned mu and sigma it will do std on train data
    standardized_data_train = standardized_vec.transform(df_tr[column_name].values.reshape(-1,1))
    print(standardized_data_train.shape)

    # with the same learned mu and sigma it will do std on cv data
    standardized_data_traincv = standardized_vec.transform(df_cv[column_name].values.reshape(-1,1))
    print(standardized_data_traincv.shape)

    # with the same learned mu and sigma it will do std on test data
    standardized_data_test = standardized_vec.transform(df_te[column_name].values.reshape(-1,1))
    print(standardized_data_test.shape)
```

```
    return standardized_data_train, standardized_data_traincv, standardized_data_test
```

In [0]:

```
from sklearn.preprocessing import Normalizer
def normalize_data(df,column_data):
    normalizer = Normalizer()
    # normalizer.fit(X_train['price'].values)
    # this will rise an error Expected 2D array, got 1D array instead:
    # array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
    # Reshape your data either using
    # array.reshape(-1, 1) if your data has a single feature
    # array.reshape(1, -1) if it contains a single sample.
    normalizer.fit(df[column_data].values.reshape(-1,1))

    data_norm = normalizer.transform(df[column_data].values.reshape(-1,1))
    print("After vectorizations")
    print(data_norm.shape)
    return data_norm
```

In [0]:

```
from sklearn.feature_extraction.text import CountVectorizer
def vectorized_data(df_train,df_cv,df_test,column_name,vocab=False):
    if(vocab):
        vectorizer = CountVectorizer(vocabulary=list(vocab.keys()), lowercase=False, binary=True)
    else:
        vectorizer = CountVectorizer(lowercase=False, binary=True)

    categories_one_hot_tr = vectorizer.fit_transform(df_train[column_name].values)
    print(vectorizer.get_feature_names())
    print("Shape of matrix after one hot encoding ",categories_one_hot_tr.shape)
    vocab_list = vectorizer.get_feature_names()

    categories_one_hot_cv = vectorizer.transform(df_cv[column_name].values)
    print(vectorizer.get_feature_names())
    print("Shape of matrix after one hot encoding ",categories_one_hot_cv.shape)

    categories_one_hot_te = vectorizer.transform(df_test[column_name].values)
    print(vectorizer.get_feature_names())
    print("Shape of matrix after one hot encoding ",categories_one_hot_te.shape)
    return categories_one_hot_tr,categories_one_hot_cv, categories_one_hot_te,vocab_list
```

In [0]:

```
from tqdm import tqdm
def textpreprocessed(df,column_name):
    # Combining all the above stundents
    global preprocessed_list

    preprocessed_list = []
    # tqdm is for printing the status bar
    for sentance in tqdm(df[column_name].values):
        sent = decontracted(sentance)
        sent = sent.replace('\\r', ' ')
        sent = sent.replace('\\'', ' ')
        sent = sent.replace('\\n', ' ')
        sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
        # https://gist.github.com/sebleier/554280
        sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
        preprocessed_list.append(sent.lower().strip())
    return preprocessed_list
```

In [0]:

```
price_standardized_tr, price_standardized_val, price_standardized_te =
standardize_data(X_train,X_val, X_test,'price')
print()
project_standardized_tr, project_standardized_val, project_standardized_te =
standardize_data(X_train,X_val, X_test,'teacher_number_of_previously_posted_projects')
```

```
(61452, 1)
(20484, 1)
(27312, 1)

(61452, 1)
(20484, 1)
(27312, 1)
```

In [0]:

```
X_train.columns
```

Out[0]:

```
Index(['Unnamed: 0', 'id', 'teacher_id', 'teacher_prefix', 'school_state',
       'Date', 'project_grade_category', 'project_title', 'project_essay_1',
       'project_essay_2', 'project_essay_3', 'project_essay_4',
       'project_resource_summary',
       'teacher_number_of_previously_posted_projects', 'price', 'quantity',
       'clean_categories', 'clean_subcategories', 'essay', 'numerical_digits',
       'project_is_approved', 'teacher_prefix0', 'teacher_prefix1',
       'school_state0', 'school_state1', 'clean_categories0',
       'clean_categories1', 'clean_subcategories0', 'clean_subcategories1',
       'project_grade_category0', 'project_grade_category1',
       'numerical_digits0'],
      dtype='object')
```

In [0]:

```
X_train.drop(['project_is_approved'], inplace=True, axis=1, errors='ignore')
```

In [0]:

```
X_train.drop(['teacher_prefix', 'school_state', 'clean_categories', 'clean_subcategories',
       'project_grade_category', 'numerical_digits','numerical_digits0'], inplace=True, axis=1, errors='ignore')
X_test.drop(['teacher_prefix', 'school_state', 'clean_categories', 'clean_subcategories',
       'project_grade_category', 'numerical_digits','numerical_digits0'], inplace=True, axis=1, errors='ignore')
X_val.drop(['teacher_prefix', 'school_state', 'clean_categories', 'clean_subcategories',
       'project_grade_category', 'numerical_digits','numerical_digits0'], inplace=True, axis=1, errors='ignore')
```

In [0]:

```
xtr = X_train[['teacher_prefix0', 'teacher_prefix1', 'school_state0', 'school_state1',
       'clean_categories0', 'clean_categories1', 'clean_subcategories0', 'clean_subcategories1',
       'project_grade_category0', 'project_grade_category1']].T.reset_index().values
xval = X_val[['teacher_prefix0', 'teacher_prefix1', 'school_state0', 'school_state1',
       'clean_categories0', 'clean_categories1', 'clean_subcategories0', 'clean_subcategories1',
       'project_grade_category0', 'project_grade_category1']].T.reset_index().values
xte = X_test[['teacher_prefix0', 'teacher_prefix1', 'school_state0', 'school_state1',
       'clean_categories0', 'clean_categories1', 'clean_subcategories0', 'clean_subcategories1',
       'project_grade_category0', 'project_grade_category1']].T.reset_index().values
```

In [0]:

```
teacher_prefix0_tr, teacher_prefix1_tr, school_state0_tr, school_state1_tr, clean_categories0_tr, c
lean_categories1_tr,\n
clean_subcategories0_tr, clean_subcategories1_tr, project_grade_category0_tr,
project_grade_category1_tr = [xtr[:,1:][i].reshape(-1,1) for i in range(10)]

teacher_prefix0_val, teacher_prefix1_val, school_state0_val, school_state1_val,
clean_categories0_val, clean_categories1_val,\n
clean_subcategories0_val, clean_subcategories1_val, project_grade_category0_val,
project_grade_category1_val = [xval[:,1:][i].reshape(-1,1) for i in range(10)]

teacher_prefix0_te, teacher_prefix1_te, school_state0_te, school_state1_te, clean_categories0_te, c
lean_categories1_te,\n
clean_subcategories0_te, clean_subcategories1_te, project_grade_category0_te,
project_grade_category1_te = [xte[:,1:][i].reshape(-1,1) for i in range(10)]
```

[Task-1] Apply both Random Forrest and GBDT on these feature sets

In [0]:

```
from scipy.sparse import hstack

f_tr = np.hstack([teacher_prefix0_tr, teacher_prefix1_tr, school_state0_tr, school_state1_tr, clean_categories0_tr, clean_categories1_tr,\n    clean_subcategories0_tr, clean_subcategories1_tr, project_grade_category0_tr,\n    project_grade_category1_tr, price_standardized_tr, project_standardized_tr])

f_cr = np.hstack([teacher_prefix0_val, teacher_prefix1_val, school_state0_val, school_state1_val, c\nlean_categories0_val, clean_categories1_val,\n    clean_subcategories0_val, clean_subcategories1_val, project_grade_category0_val,\n    project_grade_category1_val, price_standardized_val, project_standardized_val])

f_te = np.hstack([teacher_prefix0_te, teacher_prefix1_te, school_state0_te, school_state1_te, clean\n_categories0_te, clean_categories1_te,\n    clean_subcategories0_te, clean_subcategories1_te, project_grade_category0_te,\n    project_grade_category1_te, price_standardized_te, project_standardized_te])

def hstack_data(f1_tr, f1_cr, f1_te, f2_tr, f2_cr, f2_te, f3_tr, f3_cr, f3_te):\n    X_tr = hstack((f_tr.astype(float), f1_tr, f2_tr, f3_tr)).tocsr()\n    X_cr = hstack((f_cr.astype(float), f1_cr, f2_cr, f3_cr)).tocsr()\n    X_te = hstack((f_te.astype(float), f1_te, f2_te, f3_te)).tocsr()\n\n    return X_tr, X_cr, X_te
```

In [0]:

```
feature_list_x = ['teacher_prefix0', 'teacher_prefix1', 'school_state0', 'school_state1',\n    'clean_categories0', 'clean_categories1', 'clean_subcategories0', 'clean_subcategories1',\n    'project_grade_category0',\n    'project_grade_category1']+['price', 'teacher_number_of_previously_posted_projects']\nlen(feature_list_x)
```

Out[0]:

12

In [0]:

```
from sklearn.ensemble import RandomForestClassifier\nimport time\n\n\ndef optimal_hyp(X_tr,y_train,X_cr,y_val):\n    global df1\n    depth = [2,4,6,8,10,12]\n    n_estimators = [5, 10, 50, 75, 100, 200]\n    cols = ['depth', 'n_estimator', 'Train_AUC_Score', 'CV_AUC_Score']\n    lst = []\n    start = time.time()\n    for d in depth:\n        for n in n_estimators:\n            clf = RandomForestClassifier(max_depth = d, n_estimators = n, class_weight='balanced', random_state = 0)\n            clf.fit(X_tr, y_train)\n            tr_score = roc_auc_score(y_true=np.array(y_train), y_score=clf.predict_proba(X_tr)[:,1])\n            cv_score = roc_auc_score(y_true=np.array(y_val), y_score=clf.predict_proba(X_cr)[:,1])\n            #print(tr_score)\n            lst.append([d,n,tr_score,cv_score])\n    end = time.time()\n    minutes = float((end - start)/60)\n    print("execution time in minutes:",minutes)\n    print("execution time in hours:",int(minutes/60))\n\n    df1 = pd.DataFrame(lst, columns=cols)
```

In [0]:

```
from sklearn.metrics import roc_auc_score\nfrom sklearn.ensemble import GradientBoostingClassifier\nfrom xgboost import XGBClassifier
```

```

def optimal_hyp_gbdt(X_tr,y_train,X_cr,y_val):
    global df1
    depth = [2,4,6,8,10,12]
    n_estimators = [5, 10, 50,75,100]
    cols = ['depth', 'n_estimator', 'Train_AUC_Score', 'CV_AUC_Score']
    lst = []
    start = time.time()
    for d in depth:
        for n in n_estimators:
            clf = XGBClassifier(max_depth = d, n_estimators = n,class_weight='balanced',random_state = 0)
            clf.fit(X_tr, y_train)
            tr_score = roc_auc_score(y_true=np.array(y_train), y_score=clf.predict_proba(X_tr)[:,1])
            cv_score = roc_auc_score(y_true=np.array(y_val), y_score=clf.predict_proba(X_cr)[:,1])
            #print(tr_score)
            lst.append([d,n,tr_score,cv_score])
    end = time.time()
    minutes = float((end - start)/60)
    print("execution time in minutes:",minutes)
    print("execution time in hours:",int(minutes/60))

    df1 = pd.DataFrame(lst, columns=cols)

```

In [0]:

```

#https://datascience.stackexchange.com/questions/28493/confusion-matrix-get-items-fp-fn-tp-tn-python
on
def get_confusion_matrix_values(y_true, y_pred):
    cm = confusion_matrix(y_true, y_pred)
    df = pd.DataFrame(data=cm, index=labels, columns=labels)
    #tn, fp, fn, tp = cm.ravel()
    #print(tn,fp,fn,tp)
    #print("Confusion Matrix : ")
    plt.figure(figsize=(10,7))
    sns.heatmap(df, annot=True,fmt = "d")
    plt.title("Confusion Matrix",fontsize=20)
    plt.xlabel("Predicted Label",fontsize=15)
    plt.ylabel("True Label",fontsize=15)
    plt.show()
    TN, FP, FN, TP = cm[0][0], cm[0][1], cm[1][0], cm[1][1]
    print("True Positives :", TP)
    print("False Positives :", FP)
    print("True Negatives :", TN)
    print("False Negatives :", FN)

```

In [0]:

```

# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import confusion_matrix
import seaborn as sns
def ROC_plot(X_train,X_te,y_train,y_test):
    global model
    model = RandomForestClassifier(max_depth = int(best_d), n_estimators = int(n_est),class_weight='balanced', random_state = 0)
    model.fit(X_train, y_train)
    #pred = model.predict(X_train)
    y_train_pred = model.predict_proba(X_train)[:,1]

    #pred = model.predict(X_te)
    y_test_pred = model.predict_proba(X_te)[:,1]

    #y_train_pred = batch_predict(model, X_train)
    #y_test_pred = batch_predict(model, X_te)
    #print(len(y_train), len(y_train_pred))
    train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
    test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

    plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
    plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
    plt.legend()
    plt.xlabel("FPR")
    plt.ylabel("TPR")
    plt.title("ROC")

```

```
plt.title("ROC")
plt.grid()
plt.show()
print("=="*100)
```

In [0]:

```
# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sklearn.metrics.roc_curve
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import confusion_matrix
import seaborn as sns
def ROC_plot_gbdt(X_train,X_te,y_train,y_test):
    global model
    model = XGBClassifier(max_depth = int(best_d), n_estimators = int(n_est), class_weight='balanced',
random_state = 0)
    model.fit(X_train, y_train)
    #pred = model.predict(X_train)
    y_train_pred = model.predict_proba(X_train)[:,1]

    #pred = model.predict(X_te)
    y_test_pred = model.predict_proba(X_te)[:,1]

    #y_train_pred = batch_predict(model, X_train)
    #y_test_pred = batch_predict(model, X_te)
    #print(len(y_train), len(y_train_pred))
    train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
    test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

    plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
    plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
    plt.legend()
    plt.xlabel("FPR")
    plt.ylabel("TPR")
    plt.title("ROC")
    plt.grid()
    plt.show()
    print("=="*100)
```

In [0]:

```
# https://matplotlib.org/examples/mplot3d/2dcollections3d_demo.html
from mpl_toolkits.mplot3d import Axes3D
import numpy as np
import matplotlib.pyplot as plt

def plot_3D(x,y,z1,z2):
    fig = plt.figure(figsize=(10,5))
    ax = fig.gca(projection='3d')
    ax.scatter3D(x, y, z1, color="r", label='Train')
    ax.scatter3D(x, y, z2, color="b", label='CV')

    # Make legend, set axes limits and labels
    ax.legend()
    ax.set_xlabel('Depth', weight='bold')
    ax.set_ylabel('n_estimators', weight='bold')
    ax.set_zlabel('AUC', rotation=90, weight='bold')

    # Customize the view angle so it's easier to see that the scatter points lie
    # on the plane y=0
    ax.view_init(elev=20., azim=-35)

    plt.show()
```

In [0]:

```
'''# we are writing our own function for predict, with defined threshold
# we will pick a threshold that will give the least fpr
def predict(proba, threshold, fpr, tpr):

    t = threshold[np.argmax(fpr*(1-tpr))]

    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high
```

```

print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold", np.round(t,3))
predictions = []
for i in proba:
    if i>=t:
        predictions.append(1)
    else:
        predictions.append(0)
return predictions'

```

Out[0]:

```

# we are writing our own function for predict, with defined threshold\n# we will pick a threshold
that will give the least fpr\ndef predict(proba, threshold, fpr, tpr):\n    \n    t =\nthreshold[np.argmax(fpr*(1-tpr))]\n    \n    # (tpr*(1-fpr)) will be maximum if your fpr is very low\n    \n    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for\nthreshold", np.round(t,3))\n    predictions = []\n    for i in proba:\n        if i>=t:\n            predictions.append(1)\n        else:\n            predictions.append(0)\n    return predictions'

```

In [0]:

```

def plot_auc(df1):
    df = df1[['Train_AUC_Score', 'CV_AUC_Score']]
    df['depth'] = df1['depth'].astype(str)
    df['n_estimator'] = df1['n_estimator'].astype(str)

    ax = df.plot(xticks=df.index, rot=55, figsize=(15,5))
    ax.set_ylabel("AUC Score")
    ax.set_xlabel("Hyperparameters: (depth, n_estimator)")
    ax.set_xticklabels(df1[['depth', 'n_estimator']].apply(lambda x: ','.join(x), axis=1))
    ax

```

Set 2: categorical(instead of one hot encoding, try response coding: use probability values), numerical features + project_title(TFIDF) + preprocessed_eassay (TFIDF)

Vectorizing Categorical data

In [0]:

```

from sklearn.feature_extraction.text import TfidfVectorizer
def tfidf_vec(preprocessed_data_tr,preprocessed_data_val,preprocessed_data_te):
    global vectorizer_tfidf
    vectorizer_tfidf = TfidfVectorizer(min_df=10)
    text_tfidf_tr = vectorizer_tfidf.fit_transform(preprocessed_data_tr)
    vectorizer_tf = vectorizer_tfidf.get_feature_names()
    print("Shape of matrix after one hot encodig ",text_tfidf_tr.shape)

    text_tfidf_val = vectorizer_tfidf.transform(preprocessed_data_val)
    print("Shape of matrix after one hot encodig ",text_tfidf_val.shape)

    text_tfidf_te = vectorizer_tfidf.transform(preprocessed_data_te)
    print("Shape of matrix after one hot encodig ",text_tfidf_te.shape)
    return text_tfidf_tr,text_tfidf_val, text_tfidf_te, vectorizer_tf

```

In [0]:

```

tfidf_vec_essay_tr,tfidf_vec_essay_val,tfidf_vec_essay_te,tfidf_vec_essay_list =
tfidf_vec(textpreprocessed(X_train,'essay'),textpreprocessed(X_val,'essay'), textpreprocessed(X_te
st,'essay'))
tfidf_vec_titles_tr,tfidf_vec_titles_val,tfidf_vec_titles_te,tfidf_vec_titles_list =
tfidf_vec(textpreprocessed(X_train,'project_title'),textpreprocessed(X_val,'project_title'), textp
reprocessed(X_test,'project_title'))
tfidf_vec_resource_tr,tfidf_vec_resource_val,tfidf_vec_resource_te,tfidf_vec_resource_list =
tfidf_vec(textpreprocessed(X_train,'project_resource_summary'),textpreprocessed(X_val,'project_resc
urce_summary'), textpreprocessed(X_test,'project_resource_summary'))

```

100% |██████████| 61452/61452 [01:07<00:00, 908.08it/s]
100% |██████████| 20484/20484 [00:22<00:00, 913.88it/s]
100% |██████████| 27312/27312 [00:29<00:00, 914.80it/s]

```
Shape of matrix after one hot encodig (61452, 13147)
Shape of matrix after one hot encodig (20484, 13147)
```

```
3%| | 2121/61452 [00:00<00:02, 21208.71it/s]
```

```
Shape of matrix after one hot encodig (27312, 13147)
```

```
100%|██████████| 61452/61452 [00:02<00:00, 20982.64it/s]
100%|██████████| 20484/20484 [00:00<00:00, 21148.85it/s]
100%|██████████| 27312/27312 [00:01<00:00, 21099.06it/s]
```

```
Shape of matrix after one hot encodig (61452, 2289)
Shape of matrix after one hot encodig (20484, 2289)
```

```
1%| | 874/61452 [00:00<00:06, 8735.28it/s]
```

```
Shape of matrix after one hot encodig (27312, 2289)
```

```
100%|██████████| 61452/61452 [00:07<00:00, 8702.09it/s]
100%|██████████| 20484/20484 [00:02<00:00, 8675.37it/s]
100%|██████████| 27312/27312 [00:03<00:00, 8649.86it/s]
```

```
Shape of matrix after one hot encodig (61452, 4381)
Shape of matrix after one hot encodig (20484, 4381)
Shape of matrix after one hot encodig (27312, 4381)
```

In [0]:

```
feature_list = feature_list_x.copy()
feature_names = [*feature_list , *tfidf_vec_titles_list, *tfidf_vec_essay_list,
*tfidf_vec_resource_list]
print(len(feature_names))
```

19829

In [0]:

```
X_tr, X_cr, X_te = hstack_data(tfidf_vec_titles_tr, tfidf_vec_titles_val, tfidf_vec_titles_te, \
                                 tfidf_vec_essay_tr, tfidf_vec_essay_val, tfidf_vec_essay_te,\n                                 tfidf_vec_resource_tr, tfidf_vec_resource_val, tfidf_vec_resource_te)
```

In [0]:

```
X_tr.shape,X_cr.shape,X_te.shape
```

Out[0]:

```
((61452, 19829), (20484, 19829), (27312, 19829))
```

In [0]:

```
"""def plot_auc(df1):
    df = df1[['Train_AUC_Score','CV_AUC_Score']]
    df1['depth'] = df1['depth'].astype(str)
    df1['min_samples_split'] = df1['min_samples_split'].astype(str)

    ax = df.plot(xticks=df.index, rot=55,figsize=(15,5))
    ax.set_ylabel("AUC Score")
    ax.set_xlabel("Hyperparameters: (depth, min_samples_split)")
    ax.set_xticklabels(df1[['depth', 'min_samples_split']].apply(lambda x: ','.join(x), axis=1))
    ax""""
```

2.1 Applying Random Forests on TFIDF, SET 2


In [0]:

```
optimal_hyp(X_tr,y_train,X_cr,y_val)
```

```
execution time in minutes: 6.646934680143992
execution time in hours: 0
```


In [0]:

```
plot_auc(df1)
```



In [0]:

```
x = df1['depth'].astype(float).values
y = df1['n_estimator'].astype(float).values
z1 = df1['Train_AUC_Score'].astype(float).values
z2 = df1['CV_AUC_Score'].astype(float).values
plot_3D(x,y,z1,z2)
```



In [0]:

```
best_d,n_est,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]
print("optimal depth:",best_d,"\\n_estimator:",n_est,"\\nCV AUC score:", cv_auc_score)
```

```
optimal depth: 12
n_estimator: 200
CV AUC score: 0.7069679973821456
```

In [0]:

```
model = RandomForestClassifier(max_depth = int(best_d), n_estimators =
int(n_est), class_weight='balanced', random_state = 0)
model.fit(X_tr, y_train)
```

Out [0]:


```
RandomForestClassifier(bootstrap=True, class_weight='balanced',
                      criterion='gini', max_depth=12, max_features='auto',
                      max_leaf_nodes=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      n_estimators=200, n_jobs=None, oob_score=False, random_state=0,
                      verbose=0, warm_start=False)
```

In [0]:

```
y_pred_te = model.predict(X_te)
```


In [0]:

```
labels = ["negative", "positive"]
get_confusion_matrix_values(np.array(y_test), y_pred_te)
```



In [0]:

```
start = time.time()
ROC_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```





```
execution time in minutes: 0.8037584066390991  
execution time in hours: 0
```

In [0]:

```
# Printing roc auc score  
y_pred = model.predict_proba(X_te)[:,1]  
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out [0] :

0.6960855785238034

Top 100 Imp. features

In [0]:

```
#https://www.geeksforgeeks.org/generating-word-cloud-python/
from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt
import pandas as pd
coef = model.feature_importances_
coef_df = pd.DataFrame({'word': feature_names, 'coefficient': coef})
df = coef_df.sort_values("coefficient", ascending = False)[:100]
#print(df)
# iterate through the csv file
words_str = ' '
stopwords = set(STOPWORDS)
for val in df.word:
    #print(val)
    # typecaste each val to string
    val = str(val)

    # split the value
    tokens = val.split()


    # Converts each token into lowercase
    for i in range(len(tokens)):
        tokens[i] = tokens[i].lower()


    for words in tokens:
        words_str = words_str + words + ' '

wordcloud = WordCloud(width = 800, height = 800,
                      colormap="Blues",
                      stopwords = stopwords,
                      min_font_size = 10).generate(words_str)

# plot the WordCloud image
plt.figure(figsize = (8, 8), facecolor = None)
plt.title("Top 100 most important features")
plt.imshow(wordcloud)
plt.axis("off")
plt.tight_layout(pad = 0)
plt.show()
```

Top 100 most important features





2.2 Applying GBDT on TFIDF, SET 2


In [0]:

```
optimal_hyp_gbdt(X_tr,y_train,X_cr,y_val)
```

```
execution time in minutes: 73.79893393913905
execution time in hours: 1
```


In [0]:


```
plot_auc(df1)
```



In [0]:

```
x = df1['depth'].astype(float).values
y = df1['n_estimator'].astype(float).values
z1 = df1['Train_AUC_Score'].astype(float).values
z2 = df1['CV_AUC_Score'].astype(float).values
plot_3D(x,y,z1,z2)
```





In [0]:

```
best_d,n_est,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]
print("optimal depth:",best_d,"\\n_n_estimator:",n_est,"\\nCV AUC score:", cv_auc_score)
```

```
optimal depth: 12
n_estimator: 100
CV AUC score: 0.7339309134352062
```

In [0]:

```
model = XGBClassifier(max_depth = int(best_d), n_estimators = int(n_est), class_weight='balanced', random_state = 0)
model.fit(X_tr, y_train)
```

Out[0]:


```
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
              colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
              max_delta_step=0, max_depth=12, min_child_weight=1, missing=None,
              n_estimators=100, n_jobs=1, nthread=None,
              objective='binary:logistic', random_state=0, reg_alpha=0,
              reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
              subsample=1)
```

In [0]:

```
y_pred_te = model.predict(X_te)
```

In [0]:

```
labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_test), y_pred_te)
```




Predicted Label

```
True Positives : 22976
False Positives : 3863
True Negatives : 272
False Negatives : 201
```

In [0]:

```
start = time.time()
ROC_plot_gbdt(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
=====
execution time in minutes: 4.625129687786102
execution time in hours: 0
```

In [0]:

```
# Printing roc auc score
y_pred = model.predict_proba(X_te) [:,1]
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[0]:

```
0.7302925663440996
```

Top 100 Imp. features


In [0]:

```
coef = model.feature_importances_
coef_df = pd.DataFrame({'word': feature_names, 'coefficient': coef})
df = coef_df.sort_values("coefficient", ascending = False)[:100]
#print(df)
# iterate through the csv file
words_str = ' '
stopwords = set(STOPWORDS)
for val in df.word:
    #print(val)
    # typecaste each val to string
    val = str(val)

    # split the value
    tokens = val.split()

    # Converts each token into lowercase
    for i in range(len(tokens)):
        tokens[i] = tokens[i].lower()
```

```
for words in tokens:  
    words_str = words_str + words + ' '  
  
wordcloud = WordCloud(width = 800, height = 800,  
                      colormap="Blues",  
                      stopwords = stopwords,  
                      min_font_size = 10).generate(words_str)  
  
# plot the WordCloud image  
plt.figure(figsize = (8, 8), facecolor = None)  
plt.title("Top 100 most important features")  
plt.imshow(wordcloud)  
plt.axis("off")  
plt.tight_layout(pad = 0)  
plt.show()
```



Set 1: categorical, numerical features + project_title(BOW) + preprocessed_essay (BOW)

In [0]:

```
def bow_vec(preprocessed_data_tr,preprocessed_data_val,preprocessed_data_te):
    global vectorizer_bow
    vectorizer_bow = CountVectorizer(min_df=10)
    text_bow_tr = vectorizer_bow.fit_transform(preprocessed_data_tr)
    print("Shape of matrix after one hot encoding ",text_bow_tr.shape)
    vectorizer_list = vectorizer_bow.get_feature_names()
    text_bow_val = vectorizer_bow.transform(preprocessed_data_val)
    print("Shape of matrix after one hot encoding ",text_bow_val.shape)

    text_bow_te = vectorizer_bow.transform(preprocessed_data_te)
    print("Shape of matrix after one hot encoding ",text_bow_te.shape)
    return text_bow_tr,text_bow_val, text_bow_te, vectorizer_list
```

In [0]:

```
bow_vec_essay_tr, bow_vec_essay_val, bow_vec_essay_te, bow_vec_essay_list = bow_vec(textpreprocessed(X_train, 'essay'), textpreprocessed(X_val, 'essay'), textpreprocessed(X_test, 'essay'))
```

```
bow_vec_titles_tr, bow_vec_titles_val, bow_vec_titles_te, bow_vec_titles_list =  
bow_vec(textpreprocessed(X_train, 'project_title'), textpreprocessed(X_val, 'project_title'),  
textpreprocessed(X_test, 'project_title'))  
bow_vec_resource_tr, bow_vec_resource_val, bow_vec_resource_te, bow_vec_resource_list =  
bow_vec(textpreprocessed(X_train, 'project_resource_summary'), textpreprocessed(X_val, 'project_resour  
ce_summary'), textpreprocessed(X_test, 'project_resource_summary'))
```

```
100%|██████████| 61452/61452 [00:12<00:00, 5087.55it/s]  
100%|██████████| 20484/20484 [00:03<00:00, 5152.12it/s]  
100%|██████████| 27312/27312 [00:05<00:00, 5222.13it/s]
```

Shape of matrix after one hot encoding (61452, 13145)
Shape of matrix after one hot encoding (20484, 13145)

```
8%|█          | 4674/61452 [00:00<00:01, 46731.47it/s]
```

Shape of matrix after one hot encoding (27312, 13145)

```
100%|██████████| 61452/61452 [00:01<00:00, 50471.19it/s]  
100%|██████████| 20484/20484 [00:00<00:00, 49885.02it/s]  
100%|██████████| 27312/27312 [00:00<00:00, 52421.17it/s]
```

Shape of matrix after one hot encoding (61452, 2284)
Shape of matrix after one hot encoding (20484, 2284)

```
4%|█          | 2655/61452 [00:00<00:02, 26543.95it/s]
```

Shape of matrix after one hot encoding (27312, 2284)

```
100%|██████████| 61452/61452 [00:01<00:00, 32320.61it/s]  
100%|██████████| 20484/20484 [00:00<00:00, 33352.42it/s]  
100%|██████████| 27312/27312 [00:00<00:00, 33205.33it/s]
```

Shape of matrix after one hot encoding (61452, 4371)
Shape of matrix after one hot encoding (20484, 4371)
Shape of matrix after one hot encoding (27312, 4371)

In [0]:

```
feature_list = feature_list_x.copy()  
feature_names = [*feature_list, *bow_vec_titles_list, *bow_vec_essay_list, *bow_vec_resource_list]  
print(len(feature_names))
```

19812

In [0]:

```
X_tr, X_cr, X_te = hstack_data(bow_vec_titles_tr, bow_vec_titles_val, bow_vec_titles_te, \  
                                bow_vec_essay_tr, bow_vec_essay_val, bow_vec_essay_te,\ \  
                                bow_vec_resource_tr, bow_vec_resource_val, bow_vec_resource_te)
```

In [0]:

```
X_tr.shape, X_cr.shape, X_te.shape
```

Out[0]:

```
((61452, 19812), (20484, 19812), (27312, 19812))
```

1.1 Applying Random Forest on BOW, SET 1


In [0]:

```
optimal_hyp(X_tr, y_train, X_cr, y_val)
```

```
execution time in minutes: 4.3013646880785625
execution time in hours: 0
```

In [0]:

```
plot_auc(df1)
```




In [0]:

```
best_d,n_est,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]
print("optimal depth:",best_d,"\\n_estimator:",n_est,"\\nCV AUC score:", cv_auc_score)
```

```
optimal depth: 12
n_estimator: 200
CV AUC score: 0.7065557583784436
```

In [0]:

```
x = df1['depth'].astype(float).values
y = df1['n_estimator'].astype(float).values
z1 = df1['Train_AUC_Score'].astype(float).values
z2 = df1['CV_AUC_Score'].astype(float).values
plot_3D(x,y,z1,z2)
```



In [0]:

```
model = RandomForestClassifier(max_depth = int(best_d), n_estimators =
int(n_est), class_weight='balanced', random_state = 0)
model.fit(X_tr, y_train)
```

Out[0]:

```
RandomForestClassifier(bootstrap=True, class_weight='balanced',
criterion='gini', max_depth=12, max_features='auto',
```


```
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=200, n_jobs=None, oob_score=False, random_state=0,
verbose=0, warm_start=False)
```

In [0]:

```
y_pred_te = model.predict(X_te)
```


In [0]:

```
labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_test), y_pred_te)
```



In [0]:

```
start = time.time()
ROC_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
execution time in minutes: 0.6085020899772644  
execution time in hours: 0
```

In [0]:

```
# Printing roc auc score  
y_pred = model.predict_proba(X_te)[:,1]  
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[0] :

0.696862424434765

In [0]:

```
coef = model.feature_importances_
coef_df = pd.DataFrame({'word': feature_names, 'coefficient': coef})
df = coef_df.sort_values("coefficient", ascending = False)[:100]
#print(df)
# iterate through the csv file
words_str = ' '
stopwords = set(STOPWORDS)
for val in df.word:
    #print(val)
    # typecaste each val to string
    val = str(val)


    # split the value
    tokens = val.split()


    # Converts each token into lowercase
    for i in range(len(tokens)):
        tokens[i] = tokens[i].lower()

    for words in tokens:
        words_str = words_str + words + ' '

wordcloud = WordCloud(width = 800, height = 800,
                      colormap="Blues",
                      stopwords = stopwords,
                      min_font_size = 10).generate(words_str)

# plot the WordCloud image
plt.figure(figsize = (8, 8), facecolor = None)
plt.title("Top 100 most important features")
plt.imshow(wordcloud)
plt.axis("off")
plt.tight_layout(pad = 0)
plt.show()
```





1.2 Applying GBDT on BOW, SET 1

In [0]:

```
optimal_hyp_gbdt(X_tr,y_train,X_cr,y_val)
```

```
execution time in minutes: 31.82730416059494
execution time in hours: 0
```

In [0]:

```
plot_auc(df1)
```




In [0]:


```
best_d,n_est,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]
print("optimal depth:",best_d,"\\n_estimator:",n_est,"\\nCV AUC score:", cv_auc_score)
```

```
optimal depth: 10
n_estimator: 100
CV AUC score: 0.7331131121139494
```

In [0]:

```
x = df1['depth'].astype(float).values
y = df1['n_estimator'].astype(float).values
z1 = df1['Train_AUC_Score'].astype(float).values
z2 = df1['CV_AUC_Score'].astype(float).values
plot_3D(x,y,z1,z2)
```





In [0]:

```
model = XGBClassifier(max_depth = int(best_d), n_estimators = int(n_est), class_weight='balanced', random_state = 0)
model.fit(X_tr, y_train)
```

Out[0]:


```
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
  colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
  max_delta_step=0, max_depth=10, min_child_weight=1, missing=None,
  n_estimators=100, n_jobs=1, nthread=None,
  objective='binary:logistic', random_state=0, reg_alpha=0,
  reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
  subsample=1)
```

In [0]:

```
y_pred_te = model.predict(X_te)
```

In [0]:

```
labels = ["negative", "positive"]
get_confusion_matrix_values(np.array(y_test), y_pred_te)
```




```
True Positives : 23025
False Positives : 3947
True Negatives : 188
False Negatives : 152
```

In [0]:

```

start = time.time()
ROC_plot_gbdt(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))

```



```

=====
execution time in minutes: 2.050114325682322
execution time in hours: 0

```

In [0]:

```

# Printing roc auc score
y_pred = model.predict_proba(X_te) [:,1]
roc_auc_score(y_true=y_test, y_score=y_pred)

```

Out[0]:

0.7292631715583022

Top 100 Imp. features

In [0]:

```

coef = model.feature_importances_
coef_df = pd.DataFrame({'word': feature_names, 'coefficient': coef})
df = coef_df.sort_values("coefficient", ascending = False)[:100]
#print(df)
# iterate through the csv file
words_str = ' '
stopwords = set(STOPWORDS)
for val in df.word:
    #print(val)
    # typecaste each val to string
    val = str(val)

    # split the value
    tokens = val.split()

    # Converts each token into lowercase
    for i in range(len(tokens)):
        tokens[i] = tokens[i].lower()

    for words in tokens:
        words_str = words_str + words + ' '

wordcloud = WordCloud(width = 800, height = 800,
                      colormap="Blues",
                      stopwords = stopwords,
                      min_font_size = 10).generate(words_str)


# plot the WordCloud image
plt.figure(figsize = (8, 8), facecolor = None)

```

```

plt.title("Top 100 most important features")
plt.imshow(wordcloud)
plt.axis("off")
plt.tight_layout(pad = 0)
plt.show()

```



Set 3: categorical(instead of one hot encoding, try response coding: use probability values), numerical features + project_title(AVG W2V)+ preprocessed_eassay (AVG W2V)

In [0]:

```
#%cd -
```

In [0]:

```

#%cd Assignments_DonorsChoose_2018
# stronging variables into pickle files python: http://www.jessicayung.com/how-to-use-pickle-to-sa
ve-and-load-variables-in-python/
# make sure you have the glove_vectors file
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())

```

In [0]:

```

def avg_w2v(preprocessed_list):
    # average Word2Vec
    preprocessed_list = preprocessed_list[:]
    # compute average word2vec for each review.
    avg_w2v_vectors_list = [] # the avg-w2v for each sentence/review is stored in this list
    for sentence in tqdm(preprocessed_list): # for each review/sentence
        vector = np.zeros(30) # as word vectors are of zero length
        cnt_words = 0 # num of words with a valid vector in the sentence/review
        for word in sentence.split(): # for each word in a review/sentence
            if word in glove_words:
                vector += model[word] [:30]

```

```

        cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_vectors_list.append(vector)

print(len(avg_w2v_vectors_list))
print(len(avg_w2v_vectors_list[0]))
return avg_w2v_vectors_list

```

In [0]:

```

avgw2v_vec_essay_tr = avg_w2v(textpreprocessed(X_train, 'essay'))
avgw2v_vec_essay_te = avg_w2v(textpreprocessed(X_test, 'essay'))
avgw2v_vec_essay_val = avg_w2v(textpreprocessed(X_val, 'essay'))

```

```

100%|██████████| 61452/61452 [00:11<00:00, 5138.61it/s]
100%|██████████| 61452/61452 [00:15<00:00, 3840.77it/s]
 2%|          | 518/27312 [00:00<00:05, 5172.31it/s]

```

61452
30

```

100%|██████████| 27312/27312 [00:05<00:00, 5196.33it/s]
100%|██████████| 27312/27312 [00:07<00:00, 3897.92it/s]
 2%|          | 506/20484 [00:00<00:03, 5056.32it/s]

```

27312
30

```

100%|██████████| 20484/20484 [00:03<00:00, 5173.12it/s]
100%|██████████| 20484/20484 [00:05<00:00, 3822.52it/s]

```

20484
30

In [0]:

```

avgw2v_vec_titles_tr = avg_w2v(textpreprocessed(X_train, 'project_title'))
avgw2v_vec_titles_te = avg_w2v(textpreprocessed(X_test, 'project_title'))
avgw2v_vec_titles_val = avg_w2v(textpreprocessed(X_val, 'project_title'))

```

```

100%|██████████| 61452/61452 [00:01<00:00, 53112.16it/s]
100%|██████████| 61452/61452 [00:00<00:00, 84982.59it/s]
 19%|██████████| 5283/27312 [00:00<00:00, 52822.37it/s]

```

61452
30

```

100%|██████████| 27312/27312 [00:00<00:00, 52867.62it/s]
100%|██████████| 27312/27312 [00:00<00:00, 86352.40it/s]
 25%|██████████| 5164/20484 [00:00<00:00, 51639.19it/s]

```

27312
30

```

100%|██████████| 20484/20484 [00:00<00:00, 52039.88it/s]
100%|██████████| 20484/20484 [00:00<00:00, 86828.78it/s]

```

20484
30

In [0]:

```

avgw2v_vec_resource_tr = avg_w2v(textpreprocessed(X_train, 'project_resource_summary'))
avgw2v_vec_resource_te = avg_w2v(textpreprocessed(X_test, 'project_resource_summary'))
avgw2v_vec_resource_val = avg_w2v(textpreprocessed(X_val, 'project_resource_summary'))

```

```
100%|██████████| 61452/61452 [00:01<00:00, 33774.37it/s]
100%|██████████| 61452/61452 [00:01<00:00, 35939.94it/s]
12%|███          | 3373/27312 [00:00<00:00, 33725.13it/s]
```

61452
30

```
100%|██████████| 27312/27312 [00:00<00:00, 34562.10it/s]
100%|██████████| 27312/27312 [00:00<00:00, 35966.30it/s]
16%|██          | 3307/20484 [00:00<00:00, 33061.20it/s]
```

27312
30

```
100%|██████████| 20484/20484 [00:00<00:00, 33381.78it/s]
100%|██████████| 20484/20484 [00:00<00:00, 35538.23it/s]
```

20484
30

In [0]:

```
from scipy.sparse import coo_matrix
X_tr, X_cr, X_te = hstack_data(coo_matrix(np.asarray(avgw2v_vec_titles_tr)), coo_matrix(np.asarray(avgw2v_vec_titles_val)), coo_matrix(np.asarray(avgw2v_vec_titles_te)), \
                                 coo_matrix(np.asarray(avgw2v_vec_essay_tr)), coo_matrix(np.asarray(avgw2v_vec_essay_val)), coo_matrix(np.asarray(avgw2v_vec_essay_te)), \
                                 coo_matrix(np.asarray(avgw2v_vec_resource_tr)), coo_matrix(np.asarray(avgw2v_vec_resource_val)), coo_matrix(np.asarray(avgw2v_vec_resource_te)))
```

In [0]:

```
X_tr.shape,X_cr.shape,X_te.shape
```

Out[0]:

```
((61452, 102), (20484, 102), (27312, 102))
```

3.1 Applying Random Forest on AVG W2V, SET 3


In [0]:

```
optimal_hyp(X_tr,y_train,X_cr,y_val)
```

```
execution time in minutes: 38.61274026632309
execution time in hours: 0
```

In [0]:

```
plot_auc(df1)
```



2.5
2.0
2.50
2.75
2.100
2.200
4.5
4.10
4.50
4.75
4.100
4.200
6.5
6.10
6.50
6.75
6.100
6.200
8.5
8.10
8.50
8.75
8.100
8.200
10.5
10.10
10.50
10.75
10.100
10.200
12.5
12.10
12.50
12.75
12.100
12.200

Hyperparameters: (depth, n_estimator)


In [0]:

```
best_d,n_est,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]
print("optimal depth:",best_d,"n_estimator:",n_est,"CV AUC score:", cv_auc_score)
```

```
optimal depth: 10
n_estimator: 200
CV AUC score: 0.6897127326111088
```

In [0]:

```
x = df1['depth'].astype(float).values
y = df1['n_estimator'].astype(float).values
z1 = df1['Train_AUC_Score'].astype(float).values
z2 = df1['CV_AUC_Score'].astype(float).values
plot_3D(x,y,z1,z2)
```



In [0]:

```
model = RandomForestClassifier(max_depth = int(best_d), n_estimators =
int(n_est), class_weight='balanced', random_state = 0)
model.fit(X_tr, y_train)
```

Out[0]:

```
RandomForestClassifier(bootstrap=True, class_weight='balanced',
criterion='gini', max_depth=10, max_features='auto',
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=200, n_jobs=None, oob_score=False, random_state=0,
verbose=0, warm_start=False)
```

In [0]:

```
y_pred_te = model.predict(X_te)
```

In [0]:

```
labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_test), y_pred_te)
```

Confusion Matrix






```
True Positives : 19336
False Positives : 2584
True Negatives : 1551
False Negatives : 3841
```

In [0]:

```
start = time.time()
ROC_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
=====
execution time in minutes: 4.620915853977204
execution time in hours: 0
```

In [0]:

```
# Printing roc auc score
y_pred = model.predict_proba(X_te)[:,1]
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[0]:

0.673994394330075

3.2 Applying GBDT on AVG W2V, SET 3


In [0]:

```
optimal_hyp_gbdt(X_tr,y_train,X_cr,y_val)

execution time in minutes: 24.606902952988943
execution time in hours: 0
```

In [0]:

```
plot_auc(df1)
```




In [0]:

```
best_d,n_est,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]
print("optimal depth:",best_d,"\\n estimator:",n_est,"\\nCV AUC score:", cv_auc_score)
```

optimal depth: 6
n_estimator: 100
CV AUC score: 0.6999189005189344

In [0]:

```
x = df1['depth'].astype(float).values
y = df1['n_estimator'].astype(float).values
z1 = df1['Train_AUC_Score'].astype(float).values
z2 = df1['CV_AUC_Score'].astype(float).values
plot_3D(x,y,z1,z2)
```



In [0]:

```
model =XGBClassifier(max_depth = int(best_d), n_estimators = int(n_est),class_weight='balanced', random_state = 0)
model.fit(X_tr, y_train)
```

Out[0]:

Out[0]:


```
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
  colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
  max_delta_step=0, max_depth=6, min_child_weight=1, missing=None,
  n_estimators=100, n_jobs=1, nthread=None,
  objective='binary:logistic', random_state=0, reg_alpha=0,
  reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
  subsample=1)
```

In [0]:

```
y_pred_te = model.predict(X_te)
labels = ["negative","positive"]
```

In [0]:

```
get_confusion_matrix_values(np.array(y_test), y_pred_te)
```



True Positives : 23091


False Positives : 4036


True Negatives : 99

False Negatives : 86

In [0]:

```
start = time.time()
ROC_plot_gbd(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```





=====

execution time in minutes: 1.344078254699707
execution time in hours: 0

In [0]:

```
# Printing roc auc score
y_pred = model.predict_proba(X_te)[:,1]
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[0]:

0.6836673078776185

Set 4: categorical, numerical features + project_title(TFIDF W2V)+ preprocessed_essay (TFIDF W2V)

In [0]:

```
%cd Assignments_DonorsChoose_2018
# stronging variables into pickle files python: http://www.jessicayung.com/how-to-use-pickle-to-save-and-load-variables-in-python/
# make sure you have the glove_vectors file
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

In [0]:

```
# average Word2Vec
# compute average word2vec for each review.
def tfidf_w2v(preprocessed_list):
    # average Word2Vec
    preprocessed_list = preprocessed_list[:]
    tfidf_model = TfidfVectorizer()
    tfidf_model.fit(preprocessed_list)
    # we are converting a dictionary with word as a key, and the idf as a value
    dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
    tfidf_words = set(tfidf_model.get_feature_names())
    tfidf_w2v_vectors_list = [] # the avg-w2v for each sentence/review is stored in this list
    for sentence in tqdm(preprocessed_list): # for each review/sentence
        vector = np.zeros(30) # as word vectors are of zero length
        tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
        for word in sentence.split(): # for each word in a review/sentence
            if (word in glove_words) and (word in tfidf_words):
                vec = model[word][:30] # getting the vector for each word
                # here we are multiplying idf value(dictionary[word]) and the tf
                value((sentence.count(word)/len(sentence.split())))
                tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the
                tfidf value for each word
                vector += (vec * tf_idf) # calculating tfidf weighted w2v
                tf_idf_weight += tf_idf
        if tf_idf_weight != 0:
            vector /= tf_idf_weight
        tfidf_w2v_vectors_list.append(vector)

    print(len(tfidf_w2v_vectors_list))
    print(len(tfidf_w2v_vectors_list[0]))
    return tfidf_w2v_vectors_list
```

In [0]:

```
tfidfw2v_vec_essay_tr = tfidf_w2v(textpreprocessed(X_train,'essay'))
tfidfw2v_vec_essay_te = tfidf_w2v(textpreprocessed(X_test,'essay'))
```

```
tfidfw2v_vec_essay_val = tfidfw2v(textpreprocessed(X_val, 'essay'))
```

```
100%|██████████| 61452/61452 [00:12<00:00, 5079.99it/s]
100%|██████████| 61452/61452 [01:44<00:00, 586.67it/s]
 2%|██████████| 503/27312 [00:00<00:05, 5020.79it/s]
```

61452
30

```
100%|██████████| 27312/27312 [00:05<00:00, 5063.94it/s]
100%|██████████| 27312/27312 [00:46<00:00, 585.68it/s]
 2%|██████████| 497/20484 [00:00<00:04, 4964.30it/s]
```

27312
30

```
100%|██████████| 20484/20484 [00:04<00:00, 5108.40it/s]
100%|██████████| 20484/20484 [00:34<00:00, 590.57it/s]
```

20484
30

In [0]:

```
tfidfw2v_vec_titles_tr = tfidfw2v(textpreprocessed(X_train, 'project_title'))
tfidfw2v_vec_titles_te = tfidfw2v(textpreprocessed(X_test, 'project_title'))
tfidfw2v_vec_titles_val = tfidfw2v(textpreprocessed(X_val, 'project_title'))
```

```
100%|██████████| 61452/61452 [00:01<00:00, 53827.96it/s]
100%|██████████| 61452/61452 [00:01<00:00, 44901.35it/s]
 19%|██████████| 5321/27312 [00:00<00:00, 53207.89it/s]
```

61452
30

```
100%|██████████| 27312/27312 [00:00<00:00, 54036.83it/s]
100%|██████████| 27312/27312 [00:00<00:00, 45235.34it/s]
 27%|██████████| 5515/20484 [00:00<00:00, 55142.43it/s]
```

27312
30

```
100%|██████████| 20484/20484 [00:00<00:00, 54667.76it/s]
100%|██████████| 20484/20484 [00:00<00:00, 45641.85it/s]
```

20484
30

In [0]:

```
tfidfw2v_vec_resource_tr = tfidfw2v(textpreprocessed(X_train, 'project_resource_summary'))
tfidfw2v_vec_resource_te = tfidfw2v(textpreprocessed(X_test, 'project_resource_summary'))
tfidfw2v_vec_resource_val = tfidfw2v(textpreprocessed(X_val, 'project_resource_summary'))
```

```
100%|██████████| 61452/61452 [00:01<00:00, 33771.55it/s]
100%|██████████| 61452/61452 [00:04<00:00, 14204.48it/s]
 13%|██████████| 3442/27312 [00:00<00:00, 34413.63it/s]
```

61452
30

```
100%|██████████| 27312/27312 [00:00<00:00, 33832.45it/s]
100%|██████████| 27312/27312 [00:01<00:00, 14263.86it/s]
 17%|██████████| 3394/20484 [00:00<00:00, 33925.31it/s]
```

27312

```
[...]
```

```
30
```

```
100%|██████████| 20484/20484 [00:00<00:00, 33588.50it/s]
100%|██████████| 20484/20484 [00:01<00:00, 14160.96it/s]
```

```
20484
```

```
30
```

```
In [0]:
```

```
X_tr, X_cr, X_te = X_tr, X_cr, X_te = hstack_data(coo_matrix(np.asarray(tfidfw2v_vec_titles_tr)), c
oo_matrix(np.asarray(tfidfw2v_vec_titles_val)), coo_matrix(np.asarray(tfidfw2v_vec_titles_te)), \
coo_matrix(np.asarray(tfidfw2v_vec_essay_tr)), coo_matrix(np.asarray(tfidfw2v_vec_essay_val)), coo_matrix(np.asarray(tfidfw2v_vec_essay_te)), \
coo_matrix(np.asarray(tfidfw2v_vec_resource_tr)), coo_matrix(np.asarray(tfidfw2v_vec_resource_val)), coo_matrix(np.asarray(tfidfw2v_vec_resource_te)) )
```

```
In [0]:
```

```
X_tr.shape, X_cr.shape, X_te.shape
```

```
Out[0]:
```

```
((61452, 102), (20484, 102), (27312, 102))
```

4.1 Applying Random Forest on TFIDF W2V, SET 4


```
In [0]:
```

```
optimal_hyp(X_tr,y_train,X_cr,y_val)
```

```
execution time in minutes: 38.001007584730786
execution time in hours: 0
```


```
In [0]:
```

```
plot_auc(df1)
```



In [0]:

```
x = df1['depth'].astype(float).values
y = df1['n_estimator'].astype(float).values
z1 = df1['Train_AUC_Score'].astype(float).values
z2 = df1['CV_AUC_Score'].astype(float).values
plot_3D(x,y,z1,z2)
```



In [0]:

```
model = RandomForestClassifier(max_depth = int(best_d), n_estimators = int(n_est), class_weight='balanced', random_state = 0)
model.fit(X_tr, y_train)
```

Out[0]:


```
RandomForestClassifier(bootstrap=True, class_weight='balanced',
                      criterion='gini', max_depth=12, max_features='auto',
                      max_leaf_nodes=None, min_impurity_decrease=0.0,
                      min_impurity_split=None, min_samples_leaf=1,
                      min_samples_split=2, min_weight_fraction_leaf=0.0,
                      n_estimators=200, n_jobs=None, oob_score=False, random_state=0,
                      verbose=0, warm_start=False)
```


In [0]:

```
y_pred_te = model.predict(X_te)
```

In [0]:

```
labels = ["negative", "positive"]
get_confusion_matrix_values(np.array(y_test), y_pred_te)
```






```
True Positives : 20475  
False Positives : 2973  
True Negatives : 1162  
False Negatives : 2702
```

In [0]:

```
start = time.time()
ROC_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
execution time in minutes: 5.961549035708109  
execution time in hours: 0
```

In [0]:

```
# Printing roc auc score  
y_pred = model.predict_proba(X_te) [:, 1]  
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[0]:

0.6688089383530216

4.2 Applying GBDT on TFIDF W2V, SET 4


In [0]:


```
optimal_hyp qbdt(X_tr,y_train,X_cr,y_val)
```

execution time in minutes: 24.36458974679311
execution time in hours: 0

In [0]:

```
plot auc(df1)
```






In [0]:

```
best_d,n_est,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]
print("optimal depth:",best_d,"\\n_n_estimator:",n_est,"\\nCV AUC score:", cv_auc_score)
```

```
optimal depth: 6
n_estimator: 100
CV AUC score: 0.7011028809826538
```

In [0]:

```
x = df1['depth'].astype(float).values
y = df1['n_estimator'].astype(float).values
z1 = df1['Train_AUC_Score'].astype(float).values
z2 = df1['CV_AUC_Score'].astype(float).values
plot_3D(x,y,z1,z2)
```



In [0]:

```
model = XGBClassifier(max_depth = int(best_d), n_estimators = int(n_est), class_weight='balanced', random_state = 0)
model.fit(X_tr, y_train)
```

Out[0]:


```
XGBClassifier(base_score=0.5, booster='gbtree', class_weight='balanced',
              colsample_bylevel=1, colsample_bytree=1, gamma=0, learning_rate=0.1,
              max_delta_step=0, max_depth=6, min_child_weight=1, missing=None,
              n_estimators=100, n_jobs=1, nthread=None,
              objective='binary:logistic', random_state=0, reg_alpha=0,
              reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
              subsample=1)
```

In [0]:

```
y_pred_te = model.predict(X_te)
labels = ["negative", "positive"]
```

```
In [0]:
```

```
get_confusion_matrix_values(np.array(y_test), y_pred_te)
```



True Positives : 23094


False Positives : 4049

True Negatives : 86

False Negatives : 83

```
In [0]:
```

```
start = time.time()
ROC_plot_gbdt(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



=====

execution time in minutes: 1.3221298853556316

execution time in hours: 0

```
In [0]:
```

```
# Printing roc auc score
```

```
y_pred = model.predict_proba(X_te)[:,1]
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[0]:

0.6817706218466281

Summary

In [198]:

```
from prettytable import PrettyTable
import sys
sys.stdout.write("\033[1;30m")

x = PrettyTable()
x.field_names = ["Vectorizer", "Optimal Depth", "n_estimator", "AUC"]

x.add_row(["BOW", 12, 200, 0.6968])
x.add_row(["TFIDF", 12, 200, 0.6960])
x.add_row(["AVGW2V", 10, 200, 0.6739])
x.add_row(["TFIDFW2V", 12, 200, 0.6688])

print("===== RF =====")
print(x)

x = PrettyTable()
x.field_names = ["Vectorizer", "Optimal Depth", "n_estimator", "AUC"]

x.add_row(["BOW", 10, 100, 0.7302])
x.add_row(["TFIDF", 10, 100, 0.7292])
x.add_row(["AVGW2V", 6, 100, 0.6836])
x.add_row(["TFIDFW2V", 6, 100, 0.6817])

print("\n\n===== GBDT =====")
print(x)
```

```
===== RF =====
+-----+-----+-----+-----+
| Vectorizer | Optimal Depth | n_estimator | AUC |
+-----+-----+-----+-----+
| BOW | 12 | 200 | 0.6968 |
| TFIDF | 12 | 200 | 0.696 |
| AVGW2V | 10 | 200 | 0.6739 |
| TFIDFW2V | 12 | 200 | 0.6688 |
+-----+-----+-----+-----+
```

```
===== GBDT =====
+-----+-----+-----+-----+
| Vectorizer | Optimal Depth | n_estimator | AUC |
+-----+-----+-----+-----+
| BOW | 10 | 100 | 0.7302 |
| TFIDF | 10 | 100 | 0.7292 |
| AVGW2V | 6 | 100 | 0.6836 |
| TFIDFW2V | 6 | 100 | 0.6817 |
+-----+-----+-----+-----+
```

Conclusion

- GBDT outperforms RF in terms of AUC.
- BOW and TFIDF vectorizer gives better AUC than the rest of vectorizers.
- For this dataset, Model works better with depth between 6-12

Procedure followed :

- 1) Do Text Preprocessing.
- 2) Random split the dataset into train data, cross validate data and test data.
- 3) Do response coding for categorical data.
- 4) Standardize the numerical data.
- 5) Vectorize the text data by fitting the train data and then transform train, val and test data.

- 6) Train the RF and GBDT and finding the optimal n_estimator and optimal depth using the Max AUC method(used for tuning hyperparameters.)
- 7) Plot 3D plot of Train and Cv AUC score with hyperparameters depth and n_estimator respectively.
- 8) After getting the optimal n_estimator and optimal depth , Train the model again and predict the values.
- 9) Plot confusion matrix.
- 10)Plot ROC curve for train and test data respectively.

Repeat from 5) to 10) for BoW, TFIDF, Avg Word2Vec and TFIDF Word2Vec vectorizer