

```
In [1]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
In [2]: %cd drive/My Drive  
  
/content/drive/My Drive
```

```
In [3]: %cd Assignments_DonorsChoose_2018  
  
/content/drive/My Drive/Assignments_DonorsChoose_2018
```

```
In [4]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

from plotly import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter
```

```
In [0]: project_data = pd.read_csv('train_data.csv')
resource_data = pd.read_csv('resources.csv')
```

```
In [0]: # how to replace elements in list python: https://stackoverflow.com/a/2582163/408407
cols = ['Date' if x=='project_submitted_datetime' else x for x in list(project_data.columns)]
#sort dataframe based on time pandas python: https://stackoverflow.com/a/49702492
project_data['Date'] = pd.to_datetime(project_data['project_submitted_datetime'])
project_data.drop('project_submitted_datetime', axis=1, inplace=True)
project_data.sort_values(by=['Date'], inplace=True)

# how to reorder columns pandas python: https://stackoverflow.com/a/13148611/408407
project_data = project_data[cols]
#project_data.head(2)
```

```
In [0]: price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset_index()
#price_data.head(2)
```

```
In [0]: project_data = pd.merge(project_data, price_data, on='id', how='left')
```

## 1.2 preprocessing of project\_subject\_categories

```
In [0]: categories = list(project_data['project_subject_categories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/2582163/408407
# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string-in-python
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python
cat_list = []
for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space
            j=j.replace('The','') # if we have the words "The" we are going to remove it
        j = j.replace(' ','') # we are placing all the ' ' (space) with ''(empty)
        temp+=j.strip()+" "# abc ".strip() will return "abc", remove the trailing space
        temp = temp.replace('&','_') # we are replacing the & value into _
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))
```

## 1.3 preprocessing of project\_subject\_subcategories

```
In [0]: sub_catogories = list(project_data['project_subject_subcategories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-list-in-python

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-list-in-python
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

sub_cat_list = []
for i in sub_catogories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','):# it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space
            j=j.replace('The','') # if we have the words "The" we are going to remove it
        j = j.replace(' ','') # we are placing all the ' '(space) with ''(empty)
        temp +=j.strip()+" "# abc ".strip() will return "abc", remove the trailing space
        temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/4034505
my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))
```

## 1.3 Text preprocessing

```
In [0]: # merge two column text dataframe:
project_data["essay"] = project_data["project_essay_1"].map(str) + \
                        project_data["project_essay_2"].map(str) + \
                        project_data["project_essay_3"].map(str) + \
                        project_data["project_essay_4"].map(str)
```

```
In [0]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"\n\t", " not", phrase)
    phrase = re.sub(r"\re", " are", phrase)
    phrase = re.sub(r"\s", " is", phrase)
    phrase = re.sub(r"\d", " would", phrase)
    phrase = re.sub(r"\ll", " will", phrase)
    phrase = re.sub(r"\t", " not", phrase)
    phrase = re.sub(r"\ve", " have", phrase)
    phrase = re.sub(r"\m", " am", phrase)
    return phrase
```

```
In [0]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'hadn', "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', "mustn", "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'won', "won't", 'wouldn', "wouldn't"]
```

```
In [0]: def find_num(text):
    if re.findall(r'\d+', text):
        return 1
    return 0

project_data['numerical_digits'] = project_data['project_resource_summary'].apply
```

```
In [15]: project_data['project_grade_category']=project_data['project_grade_category'].str
project_data['project_grade_category']=project_data['project_grade_category'].str
set(project_data['project_grade_category'])
```

Out[15]: {'Grades\_3to5', 'Grades\_6to8', 'Grades\_9to12', 'Grades\_PreKto2'}

```
In [0]: project_data['teacher_prefix']=project_data['teacher_prefix'].fillna('Mrs.')
```

```
In [0]: #project_data = project_data[:1000]
```

```
In [0]: from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import accuracy_score
from collections import Counter
import matplotlib.pyplot as plt
import numpy as np
from scipy.sparse import csr_matrix
import time

project_data_features = project_data.copy()
project_data_features.drop('project_is_approved', axis=1, inplace=True)
y=list(project_data['project_is_approved'])
X_train, X_test, y_train, y_test = train_test_split(project_data_features, y, str
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, stratify=y_tr
```

```
In [19]: X_train.shape,X_val.shape,X_test.shape
```

```
Out[19]: ((61178, 20), (26220, 20), (21850, 20))
```

```
In [0]: from sklearn.preprocessing import StandardScaler
def standardize_data(df_tr,df_cv,df_te,column_name):
    standardized_vec = StandardScaler(with_mean=False)
    # here it will learn mu and sigma
    standardized_vec.fit(df_tr[column_name].values.reshape(-1,1))

    # with the learned mu and sigma it will do std on train data
    standardized_data_train = standardized_vec.transform(df_tr[column_name].values.
print(standardized_data_train.shape)

    # with the same learned mu and sigma it will do std on cv data
    standardized_data_traincv = standardized_vec.transform(df_cv[column_name].value
print(standardized_data_traincv.shape)

    # with the same learned mu and sigma it will do std on test data
    standardized_data_test = standardized_vec.transform(df_te[column_name].values.re
print(standardized_data_test.shape)

return standardized_data_train, standardized_data_traincv, standardized_data_te
```

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
def vectorized_data(df_train,df_cv,df_test,column_name,vocab=False):
    if(vocab):
        vectorizer = CountVectorizer(vocabulary=list(vocab.keys()), lowercase=False,
    else:
        vectorizer = CountVectorizer(lowercase=False, binary=True)

    categories_one_hot_tr = vectorizer.fit_transform(df_train[column_name].values)
    print(vectorizer.get_feature_names())
    print("Shape of matrix after one hot encoding ",categories_one_hot_tr.shape)
    vocab_list = vectorizer.get_feature_names()

    categories_one_hot_cv = vectorizer.transform(df_cv[column_name].values)
    print(vectorizer.get_feature_names())
    print("Shape of matrix after one hot encoding ",categories_one_hot_cv.shape)

    categories_one_hot_te = vectorizer.transform(df_test[column_name].values)
    print(vectorizer.get_feature_names())
    print("Shape of matrix after one hot encoding ",categories_one_hot_te.shape)
    return categories_one_hot_tr,categories_one_hot_cv, categories_one_hot_te,vocab
```

```
In [0]: def create_dict(df,column_name):
    my_counter = Counter()
    for word in df[column_name].values:
        my_counter.update(word.split())

    my_dict = dict(my_counter)
    sorted_dict = dict(sorted(my_dict.items(), key=lambda kv: kv[1]))
    return sorted_dict
```

```
In [0]: def num_hot_encode(df,column_name):
    one_hot_num_dig = pd.get_dummies(df[column_name].values)
    print("Shape of matrix after one hot encoding ",one_hot_num_dig.shape)
    return one_hot_num_dig
```

```
In [0]: from tqdm import tqdm
def textpreprocessed(df,column_name):
    # Combining all the above stundents
    preprocessed_list = []
    # tqdm is for printing the status bar
    for sentance in tqdm(df[column_name].values):
        sent = decontracted(sentance)
        sent = sent.replace('\\r', ' ')
        sent = sent.replace('\\'', ' ')
        sent = sent.replace('\\n', ' ')
        sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
        # https://gist.github.com/sebleier/554280
        sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
        preprocessed_list.append(sent.lower().strip())
    return preprocessed_list
```

```
In [0]: from sklearn.preprocessing import Normalizer
def normalize_data(df,column_data):
    normalizer = Normalizer()
    # normalizer.fit(X_train['price'].values)
    # this will rise an error Expected 2D array, got 1D array instead:
    # array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
    # Reshape your data either using
    # array.reshape(-1, 1) if your data has a single feature
    # array.reshape(1, -1) if it contains a single sample.
    normalizer.fit(df[column_data].values.reshape(-1,1))

    data_norm = normalizer.transform(df[column_data].values.reshape(-1,1))
    print("After vectorizations")
    print(data_norm.shape)
    return data_norm
```

```
In [26]: price_standardized_tr, price_standardized_val, price_standardized_te = standardize(
print()
project_standardized_tr, project_standardized_val, project_standardized_te = standardize(
    ◀ ▶
```

(61178, 1)  
(26220, 1)  
(21850, 1)

(61178, 1)  
(26220, 1)  
(21850, 1)

```
In [27]: cat_one_hot_tr,cat_one_hot_val,cat_one_hot_te,cat_one_hot_list = vectorized_data()
cat_sub_one_hot_tr,cat_sub_one_hot_val,cat_sub_one_hot_te,cat_sub_one_hot_list = 'S
    ◀ ▶
```

['Warmth', 'Care\_Hunger', 'History\_Civics', 'Music\_Arts', 'AppliedLearning', 'SpecialNeeds', 'Health\_Sports', 'Math\_Science', 'Literacy\_Language']  
Shape of matrix after one hot encoding (61178, 9)  
['Warmth', 'Care\_Hunger', 'History\_Civics', 'Music\_Arts', 'AppliedLearning', 'SpecialNeeds', 'Health\_Sports', 'Math\_Science', 'Literacy\_Language']  
Shape of matrix after one hot encoding (26220, 9)  
['Warmth', 'Care\_Hunger', 'History\_Civics', 'Music\_Arts', 'AppliedLearning', 'SpecialNeeds', 'Health\_Sports', 'Math\_Science', 'Literacy\_Language']  
Shape of matrix after one hot encoding (21850, 9)  
['Warmth', 'Care\_Hunger', 'History\_Civics', 'Music\_Arts', 'AppliedLearning', 'SpecialNeeds', 'Health\_Sports', 'Math\_Science', 'Literacy\_Language']  
Shape of matrix after one hot encoding (61178, 9)  
['Warmth', 'Care\_Hunger', 'History\_Civics', 'Music\_Arts', 'AppliedLearning', 'SpecialNeeds', 'Health\_Sports', 'Math\_Science', 'Literacy\_Language']  
Shape of matrix after one hot encoding (26220, 9)  
['Warmth', 'Care\_Hunger', 'History\_Civics', 'Music\_Arts', 'AppliedLearning', 'SpecialNeeds', 'Health\_Sports', 'Math\_Science', 'Literacy\_Language']  
Shape of matrix after one hot encoding (21850, 9)

```
In [28]: school_state_dict = create_dict(X_train,'school_state')
teacher_prefix_dict = create_dict(X_train,'teacher_prefix')

state_one_hot_tr,state_one_hot_val,state_one_hot_te,state_one_hot_list = vectorize
teacher_one_hot_tr,teacher_one_hot_val,teacher_one_hot_te,teacher_one_hot_list = '



[ 'VT', 'WY', 'ND', 'MT', 'RI', 'SD', 'NE', 'AK', 'DE', 'NH', 'WV', 'HI', 'ME',
'NM', 'DC', 'KS', 'ID', 'IA', 'AR', 'CO', 'OR', 'MN', 'KY', 'MS', 'NV', 'MD',
'CT', 'TN', 'AL', 'UT', 'WI', 'VA', 'AZ', 'NJ', 'WA', 'OK', 'MA', 'OH', 'LA',
'MO', 'IN', 'PA', 'MI', 'SC', 'GA', 'IL', 'NC', 'FL', 'NY', 'TX', 'CA']
Shape of matrix after one hot encoding (61178, 51)
[ 'VT', 'WY', 'ND', 'MT', 'RI', 'SD', 'NE', 'AK', 'DE', 'NH', 'WV', 'HI', 'ME',
'NM', 'DC', 'KS', 'ID', 'IA', 'AR', 'CO', 'OR', 'MN', 'KY', 'MS', 'NV', 'MD',
'CT', 'TN', 'AL', 'UT', 'WI', 'VA', 'AZ', 'NJ', 'WA', 'OK', 'MA', 'OH', 'LA',
'MO', 'IN', 'PA', 'MI', 'SC', 'GA', 'IL', 'NC', 'FL', 'NY', 'TX', 'CA']
Shape of matrix after one hot encoding (26220, 51)
[ 'VT', 'WY', 'ND', 'MT', 'RI', 'SD', 'NE', 'AK', 'DE', 'NH', 'WV', 'HI', 'ME',
'NM', 'DC', 'KS', 'ID', 'IA', 'AR', 'CO', 'OR', 'MN', 'KY', 'MS', 'NV', 'MD',
'CT', 'TN', 'AL', 'UT', 'WI', 'VA', 'AZ', 'NJ', 'WA', 'OK', 'MA', 'OH', 'LA',
'MO', 'IN', 'PA', 'MI', 'SC', 'GA', 'IL', 'NC', 'FL', 'NY', 'TX', 'CA']
Shape of matrix after one hot encoding (21850, 51)
['Dr.', 'Teacher', 'Mr.', 'Ms.', 'Mrs.']
Shape of matrix after one hot encoding (61178, 5)
['Dr.', 'Teacher', 'Mr.', 'Ms.', 'Mrs.']
Shape of matrix after one hot encoding (26220, 5)
['Dr.', 'Teacher', 'Mr.', 'Ms.', 'Mrs.']
Shape of matrix after one hot encoding (21850, 5)
```

```
In [29]: grade_dict = create_dict(X_train,'project_grade_category')
grade_one_hot_tr,grade_one_hot_val,grade_one_hot_te,grade_one_hot_list = vectorize
.



[ 'Grades_9to12', 'Grades_6to8', 'Grades_3to5', 'Grades_PreKto2']
Shape of matrix after one hot encoding (61178, 4)
[ 'Grades_9to12', 'Grades_6to8', 'Grades_3to5', 'Grades_PreKto2']
Shape of matrix after one hot encoding (26220, 4)
[ 'Grades_9to12', 'Grades_6to8', 'Grades_3to5', 'Grades_PreKto2']
Shape of matrix after one hot encoding (21850, 4)
```

```
In [30]: one_hot_num_dig_tr = num_hot_encode(X_train,'numerical_digits')
one_hot_num_dig_te = num_hot_encode(X_test,'numerical_digits')
one_hot_num_dig_val = num_hot_encode(X_val,'numerical_digits')
```

```
Shape of matrix after one hot encoding (61178, 2)
Shape of matrix after one hot encoding (21850, 2)
Shape of matrix after one hot encoding (26220, 2)
```

## [Task-1] Apply Support Vector Machines(SGDClassifier with hinge loss: Linear SVM) on these feature sets

```
In [0]: from scipy.sparse import hstack

f_tr = hstack((one_hot_num_dig_tr,grade_one_hot_tr,state_one_hot_tr,cat_one_hot_t
f_cr = hstack((one_hot_num_dig_val,grade_one_hot_val,state_one_hot_val,cat_one_ho
f_te = hstack((one_hot_num_dig_te,grade_one_hot_te,state_one_hot_te,cat_one_hot_te

def hstack_data(f1_tr, f1_cr, f1_te, f2_tr, f2_cr, f2_te,f3_tr,f3_cr,f3_te):
    X_tr = hstack((f_tr, f1_tr, f2_tr,f3_tr)).tocsr()
    X_cr = hstack((f_cr, f1_cr, f2_cr,f3_cr)).tocsr()
    X_te = hstack((f_te, f1_te, f2_te,f3_te)).tocsr()
    return X_tr,X_cr,X_te
```

```
In [32]: feature_list_x = ['dig_0','dig_1']+grade_one_hot_list+state_one_hot_list+cat_one_
len(feature_list_x)
```

Out[32]: 82

```
In [0]: from sklearn.metrics import roc_auc_score
from sklearn.linear_model import SGDClassifier
import time
def optimal_hyp(X_tr,y_train,X_cr,y_val):
    global df1
    global clf
    penalties = ['l1', 'l2']
    C = list(map(lambda x: 10 ** x, np.arange(-5, 3,dtype=float)))
    cols = ['C', 'Penalty', 'Train_AUC_Score', 'CV_AUC_Score']
    lst = []
    start = time.time()
    for c in C:
        for p in penalties:
            clf = SGDClassifier( class_weight='balanced', alpha=c, penalty=p, loss='h
            clf.fit(X_tr, y_train)
            tr_score = roc_auc_score(y_true=np.array(y_train), y_score=clf.decision_f
            cv_score = roc_auc_score(y_true=np.array(y_val), y_score=clf.decision_fun
            #print(tr_score)
            lst.append([c,p,tr_score,cv_score])
    end = time.time()
    minutes = float((end - start)/60)
    print("execution time in minutes:",minutes)
    print("execution time in hours:",int(minutes/60))

    df1 = pd.DataFrame(lst, columns=cols)
```

```
In [0]: #https://datascience.stackexchange.com/questions/28493/confusion-matrix-get-items
def get_confusion_matrix_values(y_true, y_pred):
    cm = confusion_matrix(y_true, y_pred)
    df = pd.DataFrame(data=cm, index=labels, columns=labels)
    #tn, fp, fn, tp = cm.ravel()
    #print(tn,fp,fn,tp)
    #print("Confusion Matrix : ")
    plt.figure(figsize=(10,7))
    sns.heatmap(df, annot=True, fmt = "d")
    plt.title("Confusion Matrix", fontsize=20)
    plt.xlabel("Predicted Label", fontsize=15)
    plt.ylabel("True Label", fontsize=15)
    plt.show()
    TN, FP, FN, TP = cm[0][0], cm[0][1], cm[1][0], cm[1][1]
    print("True Positives :", TP)
    print("False Positives :", FP)
    print("True Negatives :", TN)
    print("False Negatives :", FN)
```

```
In [0]: #https://machinelearningmastery.com/calibrated-classification-model-in-scikit-learn/
# https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import confusion_matrix
import seaborn as sns
def error_plot(X_train,X_te,y_train,y_test):
    global y_train_pred1
    model = SGDClassifier(class_weight='balanced', alpha=float(best_c), penalty=p,
    model.fit(X_train, y_train)
    #pred = model.predict(X_train)
    y_train_pred = model.decision_function(X_train)
    y_train_pred1 = model.predict(X_train)

    #pred = model.predict(X_te)
    y_test_pred = model.decision_function(X_te)

    #y_train_pred = batch_predict(model, X_train)
    #y_test_pred = batch_predict(model, X_te)
    #print(len(y_train), len(y_train_pred))
    train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
    test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)

    plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
    plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
    plt.legend()
    plt.xlabel("FPR")
    plt.ylabel("TPR")
    plt.title("ROC Curve")
    plt.grid()
    plt.show()
    print("=*100")
```

In [36]:

```
'''# we are writing our own function for predict, with defined threshold
# we will pick a threshold that will give the least fpr
def predict(proba, threshold, fpr, tpr):

    t = threshold[np.argmax(fpr*(1-tpr))]

    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high

    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold",
predictions = []
for i in proba:
    if i>=t:
        predictions.append(1)
    else:
        predictions.append(0)
return predictions'''
```

Out[36]:

```
'# we are writing our own function for predict, with defined threshold\n# we will pick a threshold that will give the least fpr\n\ndef predict(proba, threshold, fpr, tpr):\n    t = threshold[np.argmax(fpr*(1-tpr))]\n    # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very high\n\n    print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold",\n          np.round(t,3))\n    predictions = []\n    for i in proba:\n        if i>=t:\n            predictions.append(1)\n        else:\n            predictions.append(0)\n    return predictions'
```

In [0]:

```
def plot_auc(df1):
    df = df1[['Train_AUC_Score', 'CV_AUC_Score']]
    df1['C'] = df1['C'].astype(str)
    ax = df.plot(xticks=df.index, rot=45, figsize=(15,7), style='--', ms=13)
    ax.set_xticklabels(df1[['C', 'Penalty']].apply(lambda x: ''.join(x), axis=1))
    ax.set_ylabel("AUC Score")
    ax.set_xlabel("Hyperparameters: (C, Penalty)")
    ax.set_xticklabels(df1[['C', 'Penalty']].apply(lambda x: ','.join(x), axis=1))
    ax
```

```
In [0]: # https://matplotlib.org/examples/mplot3d/2dcollections3d_demo.html
from mpl_toolkits.mplot3d import Axes3D
import numpy as np
import matplotlib.pyplot as plt

def plot_3D(x,y,z1,z2):
    fig = plt.figure(figsize=(10,5))
    ax = fig.gca(projection='3d')
    ax.scatter3D(x, y, z1, color="r", label='Train')
    ax.scatter3D(x, y, z2, color="b", label='CV')

    # Make legend, set axes limits and labels
    ax.legend()
    ax.set_xlabel('C',weight='bold')
    ax.set_ylabel('Penalty',weight='bold')
    ax.set_zlabel('AUC',rotation=90,weight='bold')

    # Customize the view angle so it's easier to see that the scatter points lie
    # on the plane y=0
    ax.view_init(elev=20., azim=-35)

plt.show()
```

```
In [0]: def plot_heatmap():
    p = df1.Penalty.nunique()
    c = df1.C.nunique()
    r1=[];s1=[]
    count = 0
    for i in range(p):
        r2=[];s2=[]
        for j in range(c):
            d1 = df1.query('C=="{}" and Penalty=="{}"'.format(str(df1['C'][count]),str(df1['Penalty'][count])))
            d2 = df1.query('C=="{}" and Penalty=="{}"'.format(str(df1['C'][count]),str(df1['Penalty'][count+1])))
            count = count+1
            r2.append(d1)
            s2.append(d2)
        r1.append(r2)
        s1.append(s2)
    ##### Train AUC #####
    df = pd.DataFrame(r1, columns=list(map(lambda x: 10 ** x, np.arange(-5, 3,dtype='int'))))
    sns.heatmap(df, annot=True,fmt='.4g')
    plt.rcParams["axes.labelweight"] = "bold"
    plt.rcParams["axes.titleweight"] = "bold"
    plt.xlabel('Regularizer', fontsize=12)
    plt.ylabel('C', rotation=0, fontsize=12)
    plt.title("Train AUC Score")
    plt.show()
    ##### CV AUC #####
    df = pd.DataFrame(s1, columns=list(map(lambda x: 10 ** x, np.arange(-5, 3,dtype='int'))))
    sns.heatmap(df, annot=True,fmt='.4g')
    plt.rcParams["axes.labelweight"] = "bold"
    plt.rcParams["axes.titleweight"] = "bold"
    plt.xlabel('Regularizer', fontsize=12)
    plt.ylabel('C', rotation=0, fontsize=12)
    plt.title("CV AUC Score")
    plt.show()
```

```
In [0]: BLUE = '34m'
def display_colored_text(color, text):
    colored_text = f"\033[{color}{text}\033[00m"
    return colored_text
```

```
In [0]: #Code Reference:https://stackoverflow.com/questions/11116697/how-to-get-most-informative-features
def show_most_informative_features(feature_names, clf, n=10):
    #feature_names = vectorizer.get_feature_names()
    coefs_with_fns = sorted(zip(clf.coef_[0], feature_names))
    #print(list(coefs_with_fns))
    print(display_colored_text(BLUE, "\t\tPositive\t\t\t Negative"))
    print(' =====')
    #print("\t\tPositive\t\tNegative")
    top = zip(coefs_with_fns[:n], coefs_with_fns[-(n + 1):-1])
    for (coef_1, fn_1), (coef_2, fn_2) in top:
        print("\t%.4f\t-%15s\t\t%.4f\t-%15s" % (coef_1, fn_1, coef_2, fn_2))
```

```
In [0]: def important_features(feature_names,classifier,n=10):
    class_labels = classifier.classes_
    feature_names = feature_names
    topn_class1 = sorted(zip(classifier.feature_log_prob_[0], feature_names),reverse=True)
    topn_class2 = sorted(zip(classifier.feature_log_prob_[1], feature_names),reverse=True)
    print(display_colored_text(BLUE,"Top 10 features for negative class"))
    for coef, feat in topn_class1:
        print(class_labels[0], coef, feat)
        print("-----")
        print(display_colored_text(BLUE,"Top 10 features for positive class"))
    for coef, feat in topn_class2:
        print(class_labels[1], coef, feat)
```

## Set 1: categorical, numerical features + project\_title(BOW) + preprocessed\_eassay (BOW)

```
In [0]: def bow_vec(preprocessed_data_tr,preprocessed_data_val,preprocessed_data_te):
    global vectorizer_bow
    vectorizer_bow = CountVectorizer(min_df=10)
    text_bow_tr = vectorizer_bow.fit_transform(preprocessed_data_tr)
    print("Shape of matrix after one hot encoding ",text_bow_tr.shape)
    vectorizer_list = vectorizer_bow.get_feature_names()
    text_bow_val = vectorizer_bow.transform(preprocessed_data_val)
    print("Shape of matrix after one hot encoding ",text_bow_val.shape)

    text_bow_te = vectorizer_bow.transform(preprocessed_data_te)
    print("Shape of matrix after one hot encoding ",text_bow_te.shape)
    return text_bow_tr,text_bow_val, text_bow_te, vectorizer_list
```

```
In [0]: bow_vec_essay_tr,bow_vec_essay_val,bow_vec_essay_te,bow_vec_essay_list = bow_vec(
bow_vec_titles_tr,bow_vec_titles_val,bow_vec_titles_te,bow_vec_titles_list = bow_
bow_vec_resource_tr,bow_vec_resource_val,bow_vec_resource_te,bow_vec_resource_li
```

100% |██████████| 61178/61178 [00:37<00:00, 1632.51it/s]  
100% |██████████| 26220/26220 [00:16<00:00, 1630.63it/s]  
100% |██████████| 21850/21850 [00:13<00:00, 1635.85it/s]

Shape of matrix after one hot encoding (61178, 13127)  
Shape of matrix after one hot encoding (26220, 13127)

6% | | 3388/61178 [00:00<00:01, 33876.56it/s]

Shape of matrix after one hot encoding (21850, 13127)

100% |██████████| 61178/61178 [00:01<00:00, 34151.32it/s]  
100% |██████████| 26220/26220 [00:00<00:00, 33840.39it/s]  
100% |██████████| 21850/21850 [00:00<00:00, 34105.25it/s]

Shape of matrix after one hot encoding (61178, 2270)  
Shape of matrix after one hot encoding (26220, 2270)

3% | | 1555/61178 [00:00<00:03, 15547.90it/s]

Shape of matrix after one hot encoding (21850, 2270)

100% |██████████| 61178/61178 [00:04<00:00, 15227.36it/s]  
100% |██████████| 26220/26220 [00:01<00:00, 15236.70it/s]  
100% |██████████| 21850/21850 [00:01<00:00, 15214.15it/s]

Shape of matrix after one hot encoding (61178, 4352)  
Shape of matrix after one hot encoding (26220, 4352)  
Shape of matrix after one hot encoding (21850, 4352)

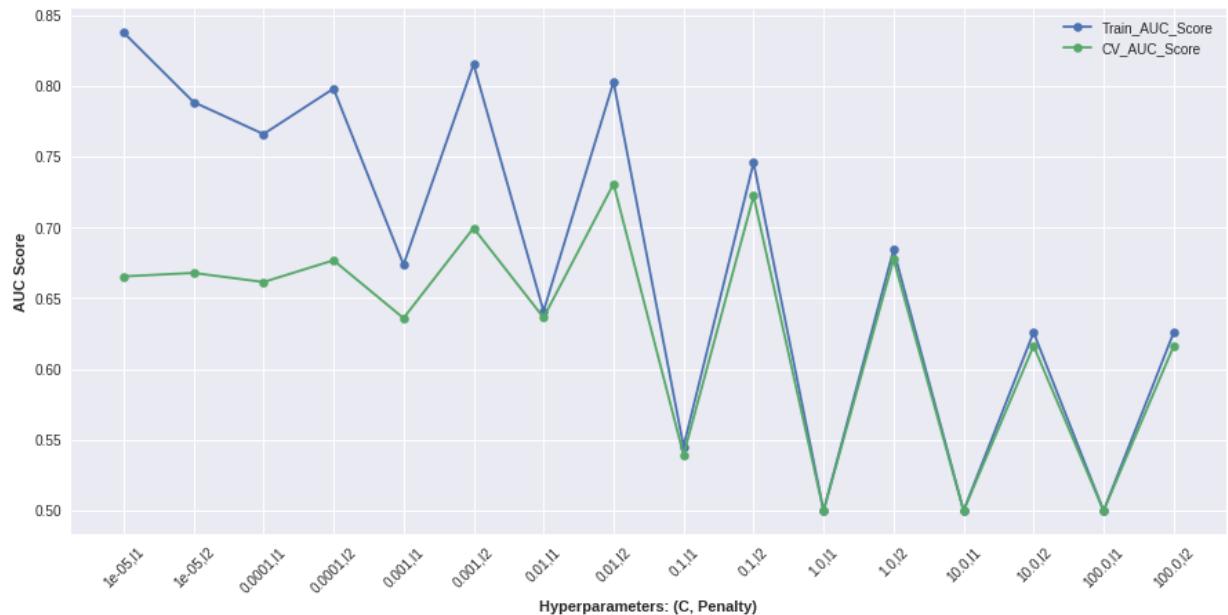
```
In [0]: X_tr, X_cr, X_te = hstack_data(bow_vec_titles_tr, bow_vec_titles_val, bow_vec_tit
bow_vec_essay_tr, bow_vec_essay_val, bow_vec_essay_te,\
bow_vec_resource_tr, bow_vec_resource_val, bow_vec_resource_te
```

```
In [0]: X_tr.shape,X_cr.shape,X_te.shape
```

```
Out[559]: ((61178, 19831), (26220, 19831), (21850, 19831))
```

```
In [0]: optimal_hyp(X_tr,y_train,X_cr,y_val)
```

```
execution time in minutes: 0.1454118847846985
execution time in hours: 0
```

In [0]: `plot_auc(df1)`

In [0]: `plot_heatmap()`



In [0]: `feature_list = feature_list_x.copy()  
feature_names = [*feature_list , *bow_vec_titles_list, *bow_vec_essay_list, *bow_vec_labels]  
print(len(feature_names))`

19831

In [0]: `best_c,p,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]  
print("optimal c:",best_c,"\\nRegularizer:",p,"\\nCV AUC score:", cv_auc_score)`

optimal c: 0.01  
Regularizer: l2  
CV AUC score: 0.730801245294767

```
In [0]: show_most_informative_features(feature_names,clf,n=10)
```

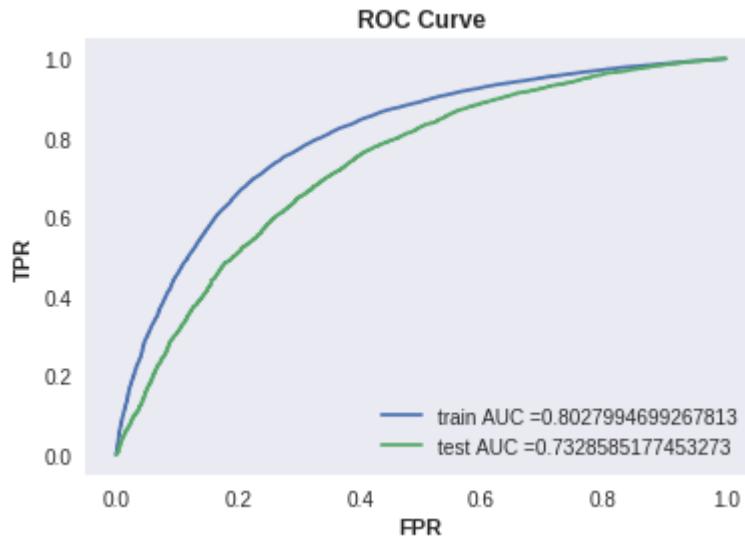
Positive	Negative
=====	
-0.0011 materials	0.0026 students
-0.0010 price	0.0013 books
-0.0007 supplies	0.0010 reading
-0.0004 supplies	0.0009 use
-0.0004 materials sted_projects	0.0009 teacher_number_of_previously_po
-0.0003 hands	0.0007 classroom
-0.0003 art	0.0007 read
-0.0003 learning	0.0006 technology
-0.0002 dig_0	0.0006 would
-0.0002 science	0.0005 day

```
In [0]: model = SGDClassifier(class_weight='balanced', alpha=float(best_c), penalty=p, lo  
model.fit(X_tr, y_train)
```

```
Out[566]: SGDClassifier(alpha=0.01, average=False, class_weight='balanced',  
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,  
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,  
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l2',  
power_t=0.5, random_state=42, shuffle=True, tol=None,  
validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [0]: y_pred = model.predict(X_te)
```

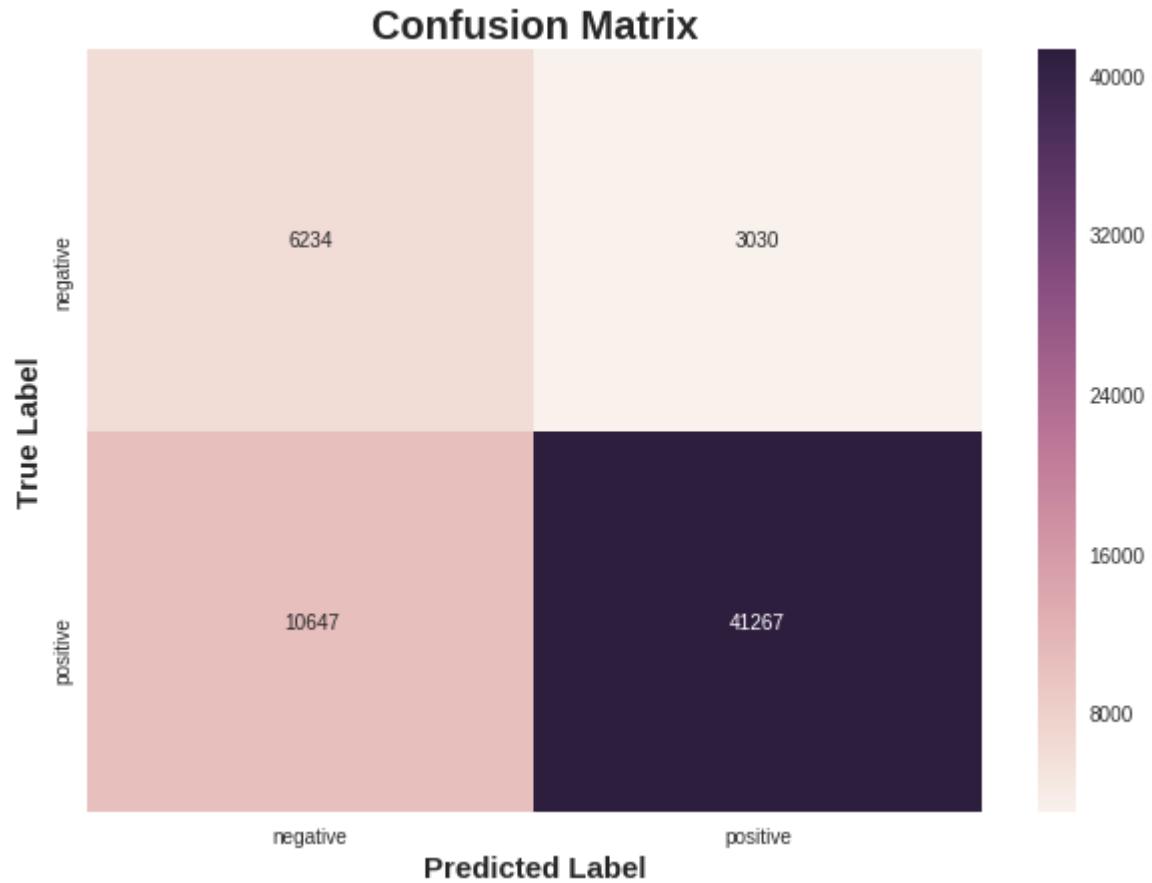
```
In [0]: start = time.time()
error_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
=====
=====
execution time in minutes: 0.011029434204101563
execution time in hours: 0
```

### confusion matrix for train data

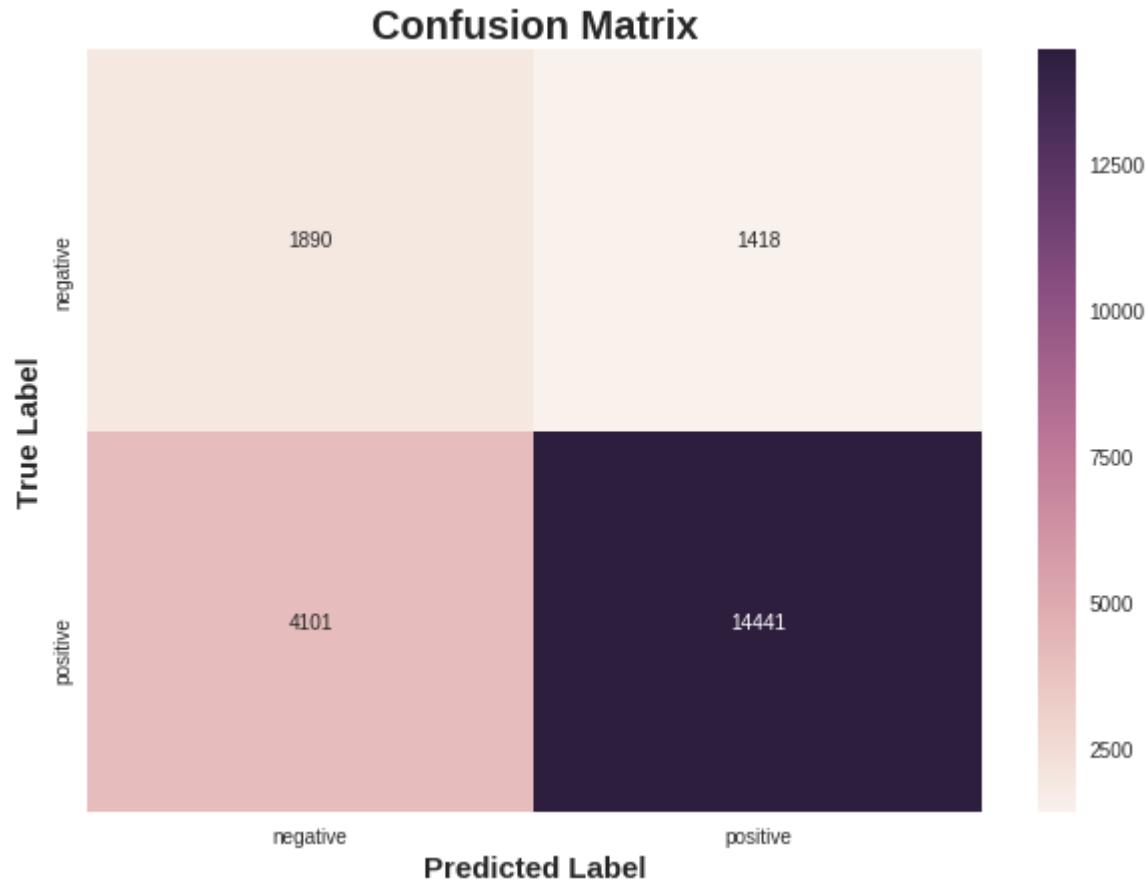
```
In [0]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_train), y_train_pred1)
```



True Positives : 41267  
False Positives : 3030  
True Negatives : 6234  
False Negatives : 10647

**confusion matrix for test data**

```
In [0]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_test), y_pred)
```



True Positives : 14441

False Positives : 1418

True Negatives : 1890

False Negatives : 4101

```
In [0]: # Printing roc auc score
y_pred = model.decision_function(x_te)
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[380]: 0.7289469920701614

**Set 2: categorical, numerical features + project\_title(TFIDF)+ preprocessed\_eassay (TFIDF)**

```
In [0]: from sklearn.feature_extraction.text import TfidfVectorizer
def tfidf_vec(preprocessed_data_tr, preprocessed_data_val, preprocessed_data_te):
    global vectorizer_tfidf
    vectorizer_tfidf = TfidfVectorizer(min_df=10)
    text_tfidf_tr = vectorizer_tfidf.fit_transform(preprocessed_data_tr)
    vectorizer_tf = vectorizer_tfidf.get_feature_names()
    print("Shape of matrix after one hot encoding ", text_tfidf_tr.shape)

    text_tfidf_val = vectorizer_tfidf.transform(preprocessed_data_val)
    print("Shape of matrix after one hot encoding ", text_tfidf_val.shape)

    text_tfidf_te = vectorizer_tfidf.transform(preprocessed_data_te)
    print("Shape of matrix after one hot encoding ", text_tfidf_te.shape)
    return text_tfidf_tr, text_tfidf_val, text_tfidf_te, vectorizer_tf
```

```
In [0]: tfidf_vec_essay_tr, tfidf_vec_essay_val, tfidf_vec_essay_te, tfidf_vec_essay_list =
tfidf_vec_titles_tr, tfidf_vec_titles_val, tfidf_vec_titles_te, tfidf_vec_titles_list =
tfidf_vec_resource_tr, tfidf_vec_resource_val, tfidf_vec_resource_te, tfidf_vec_reso
```

100% | ██████████ | 61178/61178 [00:37<00:00, 1614.04it/s]  
100% | ██████████ | 26220/26220 [00:16<00:00, 1624.80it/s]  
100% | ██████████ | 21850/21850 [00:13<00:00, 1640.09it/s]

Shape of matrix after one hot encoding (61178, 13174)  
Shape of matrix after one hot encoding (26220, 13174)

5% | | 3325/61178 [00:00<00:01, 33244.72it/s]

Shape of matrix after one hot encoding (21850, 13174)

100% | ██████████ | 61178/61178 [00:01<00:00, 34218.14it/s]  
100% | ██████████ | 26220/26220 [00:00<00:00, 34045.14it/s]  
100% | ██████████ | 21850/21850 [00:00<00:00, 34076.63it/s]

Shape of matrix after one hot encoding (61178, 2292)  
Shape of matrix after one hot encoding (26220, 2292)

3% | | 1541/61178 [00:00<00:03, 15404.84it/s]

Shape of matrix after one hot encoding (21850, 2292)

100% | ██████████ | 61178/61178 [00:04<00:00, 15292.60it/s]  
100% | ██████████ | 26220/26220 [00:01<00:00, 15165.00it/s]  
100% | ██████████ | 21850/21850 [00:01<00:00, 15150.08it/s]

Shape of matrix after one hot encoding (61178, 4369)  
Shape of matrix after one hot encoding (26220, 4369)  
Shape of matrix after one hot encoding (21850, 4369)

```
In [0]: feature_list = feature_list_x.copy()
feature_names = [*feature_list, *tfidf_vec_titles_list, *tfidf_vec_essay_list, *print(len(feature_names))]
```

19917

```
In [0]: X_tr, X_cr, X_te = hstack_data(tfidf_vec_titles_tr, tfidf_vec_titles_val, tfidf_vec_titles_te,\n    tfidf_vec_essay_tr, tfidf_vec_essay_val, tfidf_vec_essay_te,\n    tfidf_vec_resource_tr, tfidf_vec_resource_val, tfidf_vec_resou
```

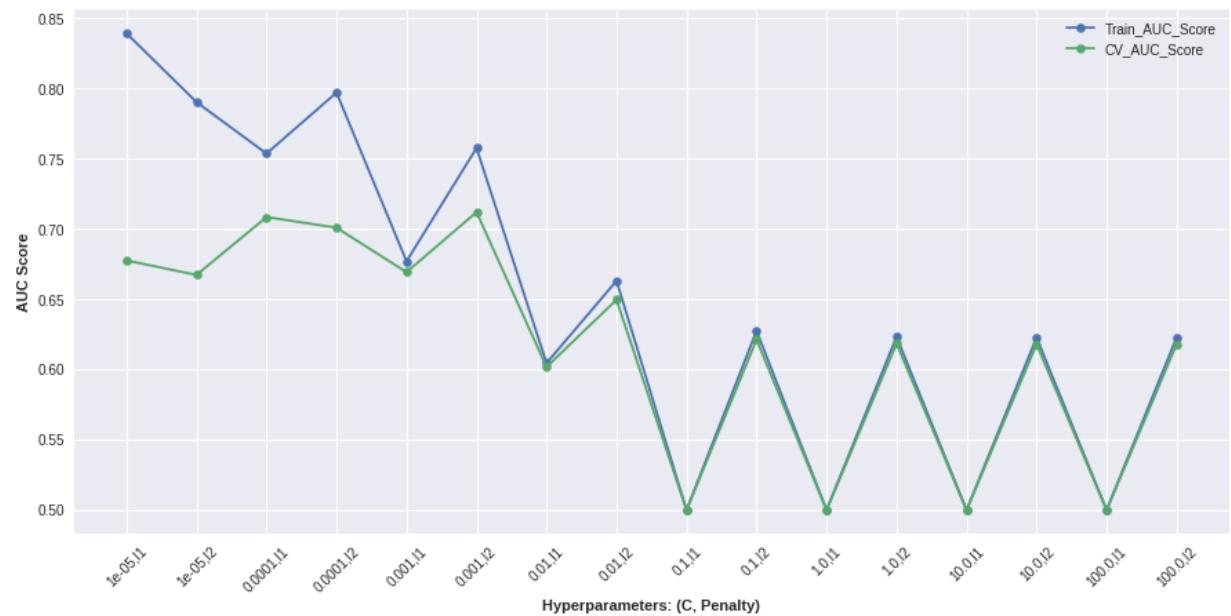
```
In [0]: X_tr.shape,X_cr.shape,X_te.shape
```

```
Out[617]: ((61178, 19917), (26220, 19917), (21850, 19917))
```

```
In [0]: optimal_hyp(X_tr,y_train,X_cr,y_val)
```

execution time in minutes: 0.14380438327789308  
 execution time in hours: 0

```
In [0]: plot_auc(df1)
```



In [0]: `plot_heatmap()`



In [0]: `best_c,p,tr_auc_score,cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]  
print("optimal c:",best_c,"\\nRegularizer:",p,"\\nCV AUC score:", cv_auc_score)`

optimal c: 0.001  
 Regularizer: 12  
 CV AUC score: 0.7120557665638355

In [0]: `show_most_informative_features(feature_names, clf, n=10)`

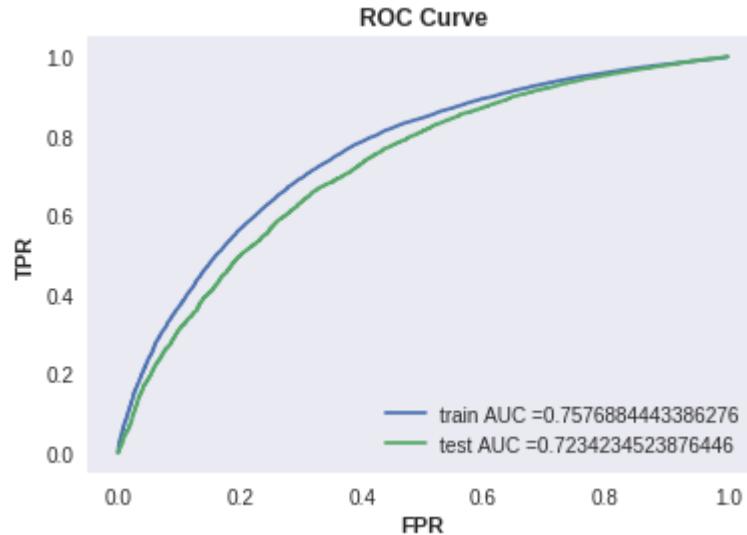
Positive	Negative
=====	
-0.0010 price	0.0009 teacher_number_of_previously_posted_projects
-0.0002 dig_0	0.0003 Literacy_Language
-0.0001 TX	0.0002 dig_1
-0.0001 Math_Science	0.0001 Grades_3to5
-0.0001 SpecialNeeds	0.0001 books
-0.0001 SpecialNeeds	0.0000 CA
-0.0001 supplies	0.0000 reading
-0.0001 materials	0.0000 books
-0.0001 AppliedLearning	0.0000 chromebooks
-0.0001 materials	0.0000 chromebooks

In [0]: `model = SGDClassifier(class_weight='balanced', alpha=float(best_c), penalty=p, loss='hinge', max_iter=None, tol=None, random_state=42, shuffle=True, verbose=0, warm_start=False)`  
`model.fit(X_tr, y_train)`

Out[623]: `SGDClassifier(alpha=0.001, average=False, class_weight='balanced', early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True, l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None, n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l2', power_t=0.5, random_state=42, shuffle=True, tol=None, validation_fraction=0.1, verbose=0, warm_start=False)`

In [0]: `y_pred = model.predict(X_te)`

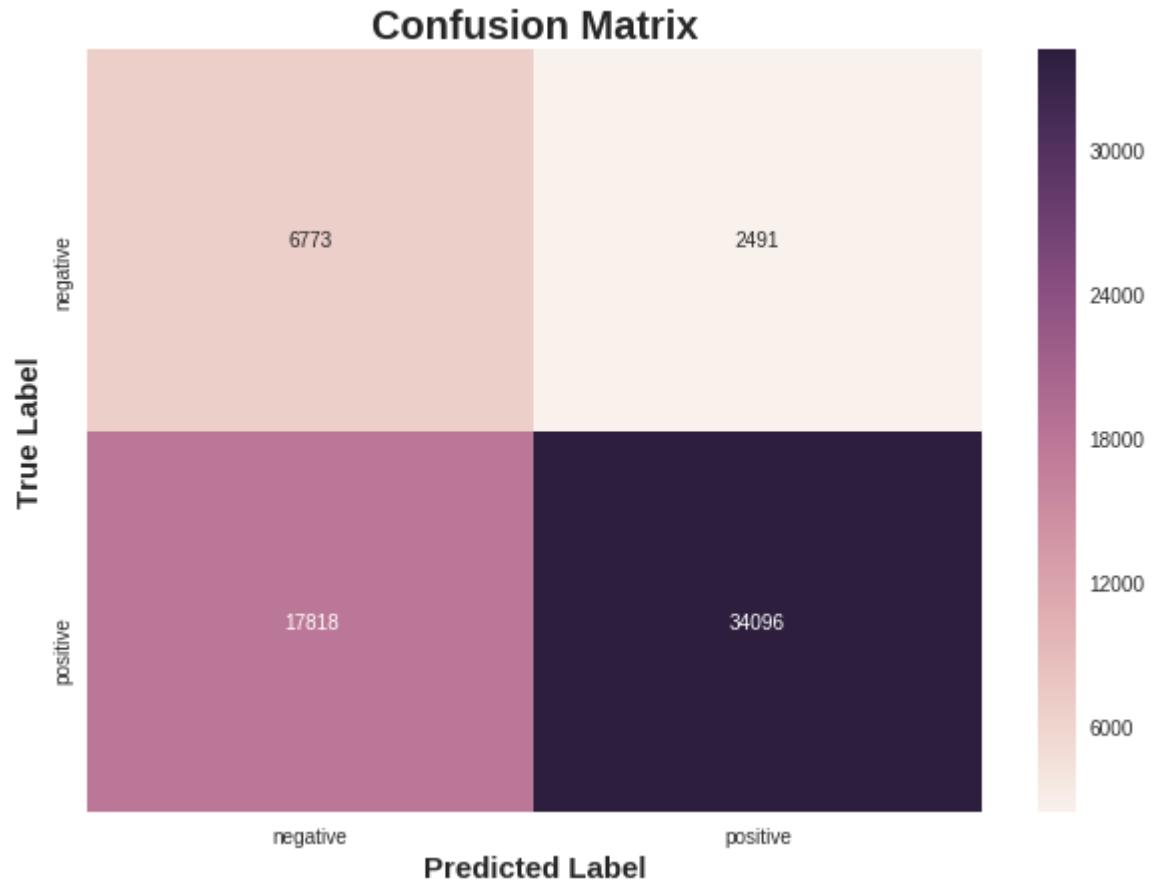
```
In [0]: start = time.time()
error_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
=====
=====
execution time in minutes: 0.012538317839304607
execution time in hours: 0
```

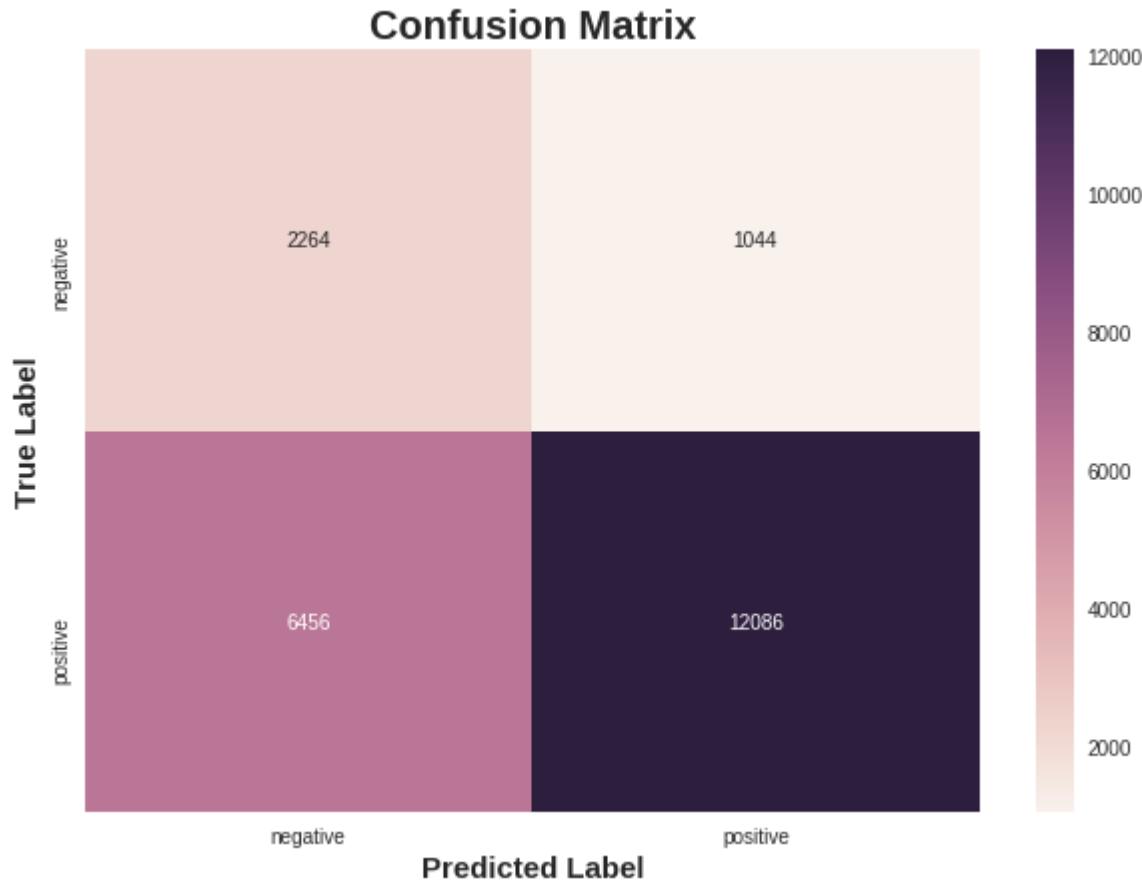
### confusion matrix for train data

```
In [0]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_train), y_train_pred1)
```



**confusion matrix for test data**

```
In [0]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_test), y_pred)
```



True Positives : 12086

False Positives : 1044

True Negatives : 2264

False Negatives : 6456

```
In [0]: # Printing roc auc score
y_pred = model.decision_function(x_te)
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[628]: 0.7234234523876446

### Set 3: categorical, numerical features + project\_title(AVG W2V)+ preprocessed\_eassay (AVG W2V)

```
In [0]: #cd -
```

```
In [0]: #cd Assignments_DonorsChoose_2018
# stronging variables into pickle files python: http://www.jessicayung.com/how-to
# make sure you have the glove_vectors file
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

```
In [0]: def avg_w2v(preprocessed_list):
    # average Word2Vec
    preprocessed_list = preprocessed_list[:]
    # compute average word2vec for each review.
    avg_w2v_vectors_list = [] # the avg-w2v for each sentence/review is stored in
    for sentence in tqdm(preprocessed_list): # for each review/sentence
        vector = np.zeros(30) # as word vectors are of zero length
        cnt_words = 0; # num of words with a valid vector in the sentence/review
        for word in sentence.split(): # for each word in a review/sentence
            if word in glove_words:
                vector += model[word][:30]
                cnt_words += 1
        if cnt_words != 0:
            vector /= cnt_words
        avg_w2v_vectors_list.append(vector)

    print(len(avg_w2v_vectors_list))
    print(len(avg_w2v_vectors_list[0]))
    return avg_w2v_vectors_list
```

```
In [772]: avgw2v_vec_essay_tr = avg_w2v(textpreprocessed(X_train,'essay'))
avgw2v_vec_essay_te = avg_w2v(textpreprocessed(X_test,'essay'))
avgw2v_vec_essay_val = avg_w2v(textpreprocessed(X_val,'essay'))
```

100% |██████████| 61178/61178 [00:37<00:00, 1623.16it/s]  
100% |██████████| 61178/61178 [00:16<00:00, 3703.48it/s]  
1% | 160/21850 [00:00<00:13, 1588.53it/s]

61178  
30

100% |██████████| 21850/21850 [00:13<00:00, 1622.09it/s]  
100% |██████████| 21850/21850 [00:05<00:00, 3701.41it/s]  
1% | 160/26220 [00:00<00:16, 1592.29it/s]

21850  
30

100% |██████████| 26220/26220 [00:16<00:00, 1621.29it/s]  
100% |██████████| 26220/26220 [00:07<00:00, 3613.27it/s]

26220  
30

```
In [773]: avgw2v_vec_titles_tr = avg_w2v(textpreprocessed(X_train,'project_title'))
avgw2v_vec_titles_te = avg_w2v(textpreprocessed(X_test,'project_title'))
avgw2v_vec_titles_val = avg_w2v(textpreprocessed(X_val,'project_title'))
```

100% | ██████████ | 61178/61178 [00:01<00:00, 34026.09it/s]  
100% | ██████████ | 61178/61178 [00:00<00:00, 86912.87it/s]  
15% | █ | 3386/21850 [00:00<00:00, 33857.93it/s]

61178

30

100% | ██████████ | 21850/21850 [00:00<00:00, 33970.71it/s]  
100% | ██████████ | 21850/21850 [00:00<00:00, 83809.75it/s]  
13% | █ | 3439/26220 [00:00<00:00, 34387.74it/s]

21850

30

100% | ██████████ | 26220/26220 [00:00<00:00, 34417.28it/s]  
100% | ██████████ | 26220/26220 [00:00<00:00, 86640.26it/s]

26220

30

```
In [774]: avgw2v_vec_resource_tr = avg_w2v(textpreprocessed(X_train,'project_resource_summary'))
avgw2v_vec_resource_te = avg_w2v(textpreprocessed(X_test,'project_resource_summary'))
avgw2v_vec_resource_val = avg_w2v(textpreprocessed(X_val,'project_resource_summary'))
```

100% | ██████████ | 61178/61178 [00:03<00:00, 15363.55it/s]  
100% | ██████████ | 61178/61178 [00:01<00:00, 35414.83it/s]  
7% | █ | 1475/21850 [00:00<00:01, 14747.10it/s]

61178

30

100% | ██████████ | 21850/21850 [00:01<00:00, 15219.00it/s]  
100% | ██████████ | 21850/21850 [00:00<00:00, 34917.44it/s]  
6% | █ | 1466/26220 [00:00<00:01, 14655.75it/s]

21850

30

100% | ██████████ | 26220/26220 [00:01<00:00, 15127.62it/s]  
100% | ██████████ | 26220/26220 [00:00<00:00, 35339.89it/s]

26220

30

```
In [0]: X_tr, X_cr, X_te = hstack_data(avgw2v_vec_titles_tr, avgw2v_vec_titles_val, avgw2v_vec_titles_te,
                                     avgw2v_vec_essay_tr, avgw2v_vec_essay_val, avgw2v_vec_essay_te,
                                     avgw2v_vec_resource_tr, avgw2v_vec_resource_val, avgw2v_vec_re
```

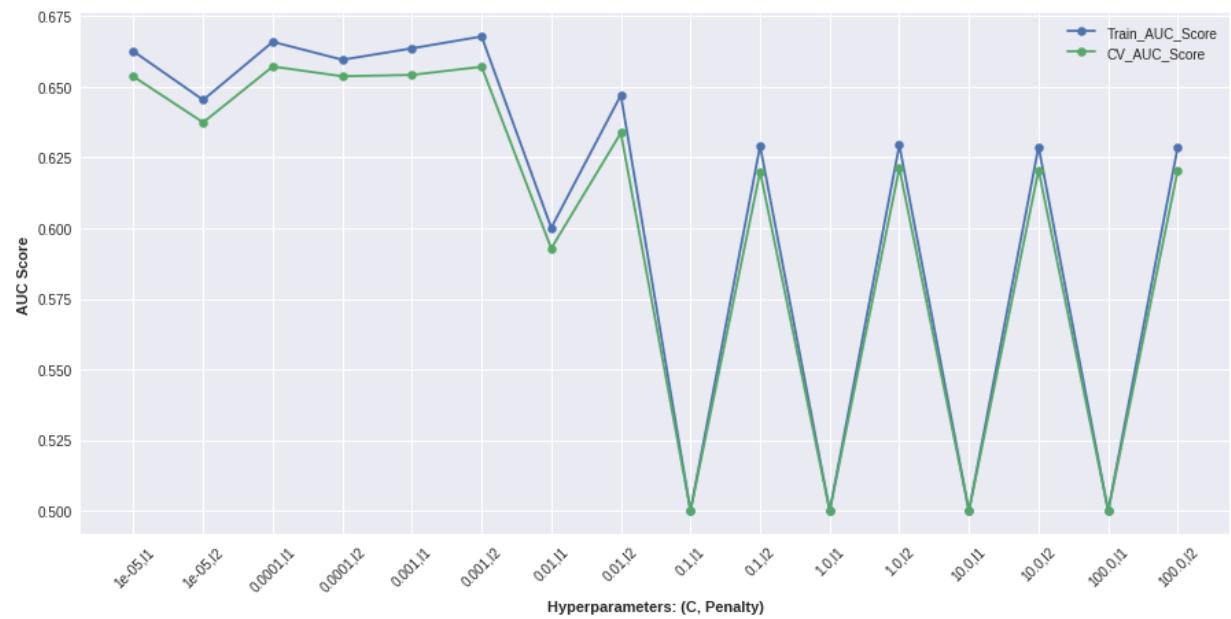
```
In [776]: X_tr.shape,X_cr.shape,X_te.shape
```

```
Out[776]: ((61178, 172), (26220, 172), (21850, 172))
```

```
In [777]: optimal_hyp(X_tr,y_train,X_cr,y_val)
```

```
execution time in minutes: 0.12700416644414267  
execution time in hours: 0
```

```
In [778]: plot_auc(df1)
```



In [779]: `plot_heatmap()`



In [780]: `best_c, p, tr_auc_score, cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]  
print("optimal c:", best_c, "\nRegularizer:", p, "\nCV AUC score:", cv_auc_score)`

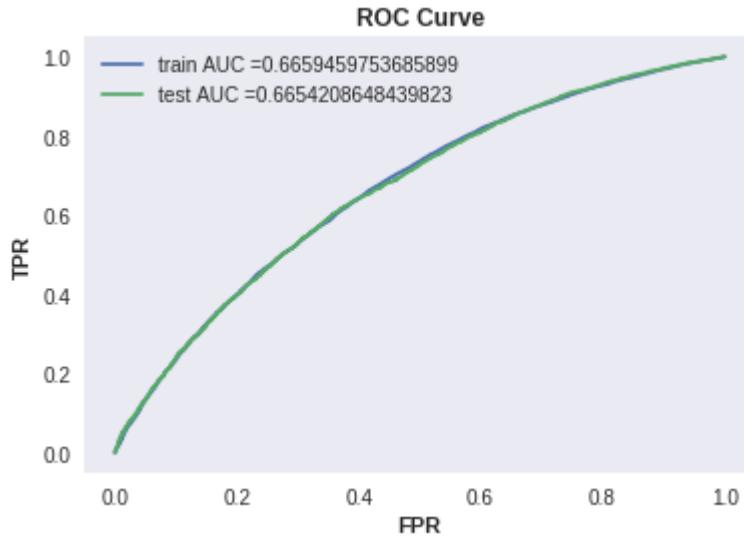
optimal c: 0.0001  
Regularizer: l1  
CV AUC score: 0.6572138001301899

In [781]: `model = SGDClassifier(class_weight='balanced', alpha=float(best_c), penalty=p, l0  
model.fit(X_tr, y_train)`

Out[781]: `SGDClassifier(alpha=0.0001, average=False, class_weight='balanced',  
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,  
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,  
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l1',  
power_t=0.5, random_state=42, shuffle=True, tol=None,  
validation_fraction=0.1, verbose=0, warm_start=False)`

In [0]: `y_pred = model.predict(X_te)`

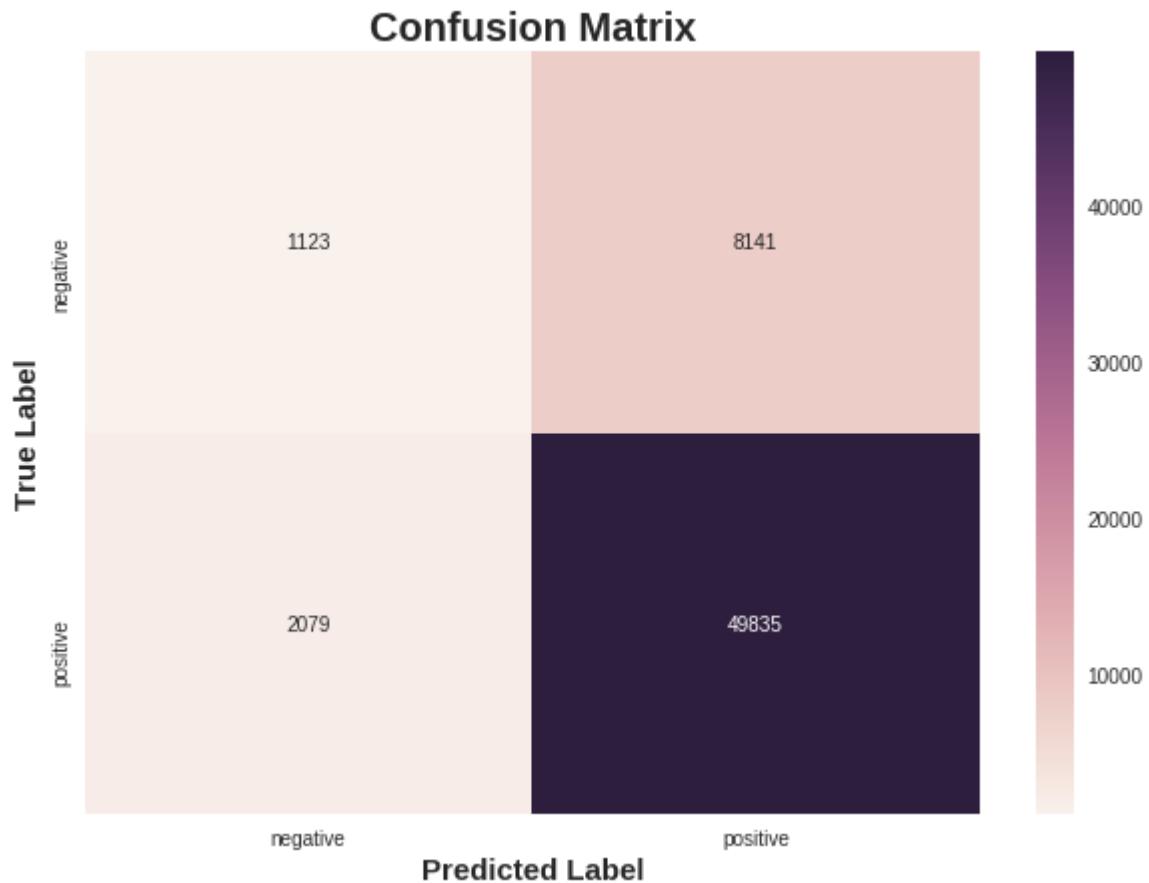
```
In [783]: start = time.time()
error_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
=====
=====
execution time in minutes: 0.01321113109588623
execution time in hours: 0
```

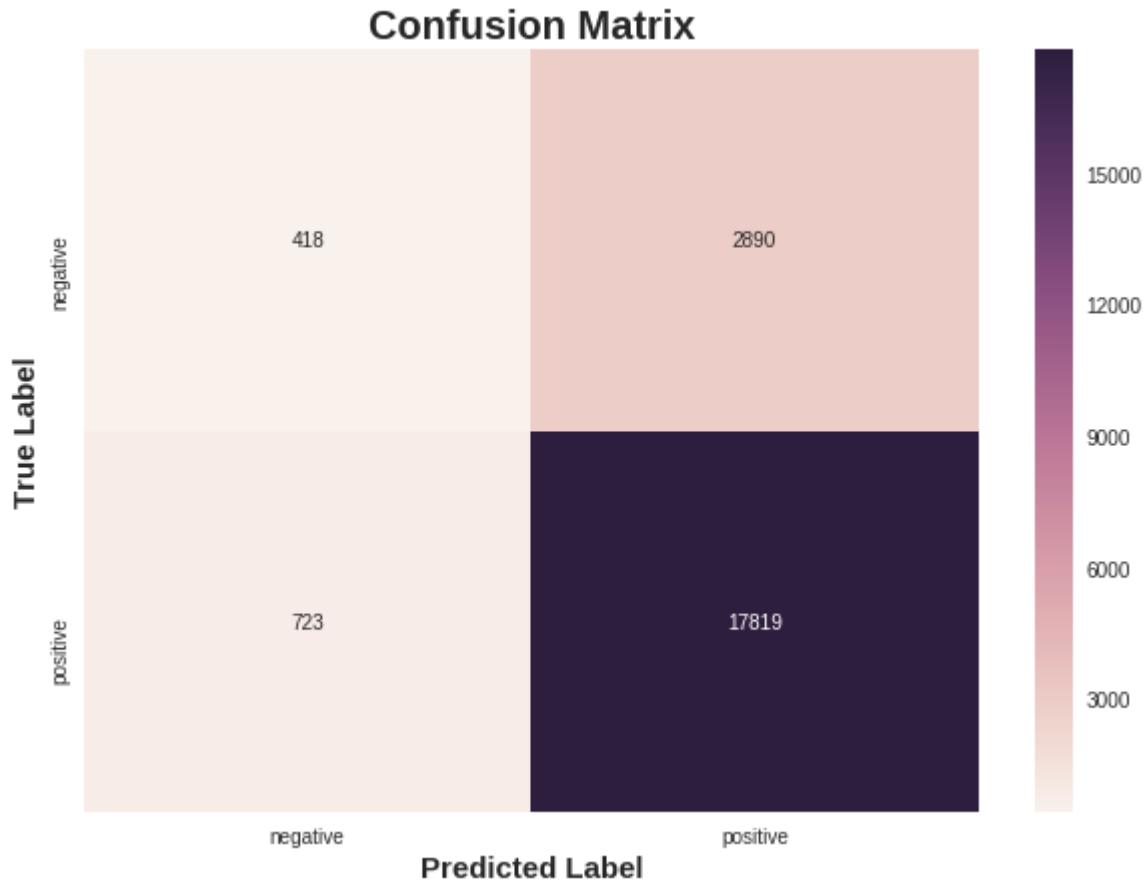
### confusion matrix for train data

```
In [784]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_train), y_train_pred1)
```



**confusion matrix for test data**

```
In [785]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_test), y_pred)
```



True Positives : 17819  
 False Positives : 2890  
 True Negatives : 418  
 False Negatives : 723

```
In [786]: # Printing roc auc score
y_pred = model.decision_function(x_te)
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[786]: 0.6654208648439823

#### Set 4: categorical, numerical features + project\_title(TFIDF W2V)+ preprocessed\_eassay (TFIDF W2V)

```
In [0]: #cd Assignments_DonorsChoose_2018
# stronging variables into pickle files python: http://www.jessicayung.com/how-to
# make sure you have the glove_vectors file
with open('glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

```
In [0]: # average Word2Vec
# compute average word2vec for each review.
def tfidf_w2v(preprocessed_list):
    # average Word2Vec
    preprocessed_list = preprocessed_list[:]
    tfidf_model = TfidfVectorizer()
    tfidf_model.fit(preprocessed_list)
    # we are converting a dictionary with word as a key, and the idf as a value
    dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
    tfidf_words = set(tfidf_model.get_feature_names())
    tfidf_w2v_vectors_list = [] # the avg-w2v for each sentence/review is stored in
    for sentence in tqdm(preprocessed_list): # for each review/sentence
        vector = np.zeros(30) # as word vectors are of zero length
        tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
        for word in sentence.split(): # for each word in a review/sentence
            if (word in glove_words) and (word in tfidf_words):
                vec = model[word][:30] # getting the vector for each word
                # here we are multiplying idf value(dictionary[word]) and the tf value
                tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split()))
                vector += (vec * tf_idf) # calculating tfidf weighted w2v
                tf_idf_weight += tf_idf
        if tf_idf_weight != 0:
            vector /= tf_idf_weight
        tfidf_w2v_vectors_list.append(vector)

    print(len(tfidf_w2v_vectors_list))
    print(len(tfidf_w2v_vectors_list[0]))
    return tfidf_w2v_vectors_list
```

```
In [789]: tfidfw2v_vec_essay_tr = tfidf_w2v(textpreprocessed(X_train,'essay'))
tfidfw2v_vec_essay_te = tfidf_w2v(textpreprocessed(X_test,'essay'))
tfidfw2v_vec_essay_val = tfidf_w2v(textpreprocessed(X_val,'essay'))
```

100% | ██████████ | 61178/61178 [00:37<00:00, 1631.66it/s]  
100% | ██████████ | 61178/61178 [01:47<00:00, 569.40it/s]  
1% | 165/21850 [00:00<00:13, 1639.50it/s]

61178

30

100% | ██████████ | 21850/21850 [00:13<00:00, 1618.38it/s]  
100% | ██████████ | 21850/21850 [00:38<00:00, 560.73it/s]  
1% | 160/26220 [00:00<00:16, 1592.66it/s]

21850

30

100% | ██████████ | 26220/26220 [00:16<00:00, 1604.35it/s]  
100% | ██████████ | 26220/26220 [00:45<00:00, 570.40it/s]

26220

30

```
In [790]: tfidfw2v_vec_titles_tr = tfidf_w2v(textpreprocessed(X_train,'project_title'))
tfidfw2v_vec_titles_te = tfidf_w2v(textpreprocessed(X_test,'project_title'))
tfidfw2v_vec_titles_val = tfidf_w2v(textpreprocessed(X_val,'project_title'))
```

100% |██████████| 61178/61178 [00:01<00:00, 34276.99it/s]  
100% |██████████| 61178/61178 [00:01<00:00, 43769.59it/s]  
16% |██| 3503/21850 [00:00<00:00, 35022.85it/s]

61178

30

100% |██████████| 21850/21850 [00:00<00:00, 34027.19it/s]  
100% |██████████| 21850/21850 [00:00<00:00, 44026.43it/s]  
13% |██| 3354/26220 [00:00<00:00, 33536.11it/s]

21850

30

100% |██████████| 26220/26220 [00:00<00:00, 34105.81it/s]  
100% |██████████| 26220/26220 [00:00<00:00, 43717.36it/s]

26220

30

```
In [791]: tfidfw2v_vec_resource_tr = tfidf_w2v(textpreprocessed(X_train,'project_resource_s')
tfidfw2v_vec_resource_te = tfidf_w2v(textpreprocessed(X_test,'project_resource_su
tfidfw2v_vec_resource_val = tfidf_w2v(textpreprocessed(X_val,'project_resource_su
```

100% |██████████| 61178/61178 [00:04<00:00, 15268.20it/s]  
100% |██████████| 61178/61178 [00:04<00:00, 13700.99it/s]  
7% |██| 1475/21850 [00:00<00:01, 14746.04it/s]

61178

30

100% |██████████| 21850/21850 [00:01<00:00, 15053.49it/s]  
100% |██████████| 21850/21850 [00:01<00:00, 13944.26it/s]  
6% |██| 1514/26220 [00:00<00:01, 15134.93it/s]

21850

30

100% |██████████| 26220/26220 [00:01<00:00, 15313.30it/s]  
100% |██████████| 26220/26220 [00:01<00:00, 14136.70it/s]

26220

30

```
In [0]: X_tr, X_cr, X_te = hstack_data(tfidfw2v_vec_titles_tr, tfidfw2v_vec_titles_val, t
tfidfw2v_vec_essay_tr, tfidfw2v_vec_essay_val, tfidfw2v_vec_es
tfidfw2v_vec_resource_tr, tfidfw2v_vec_resource_val, tfidfw2v_
```

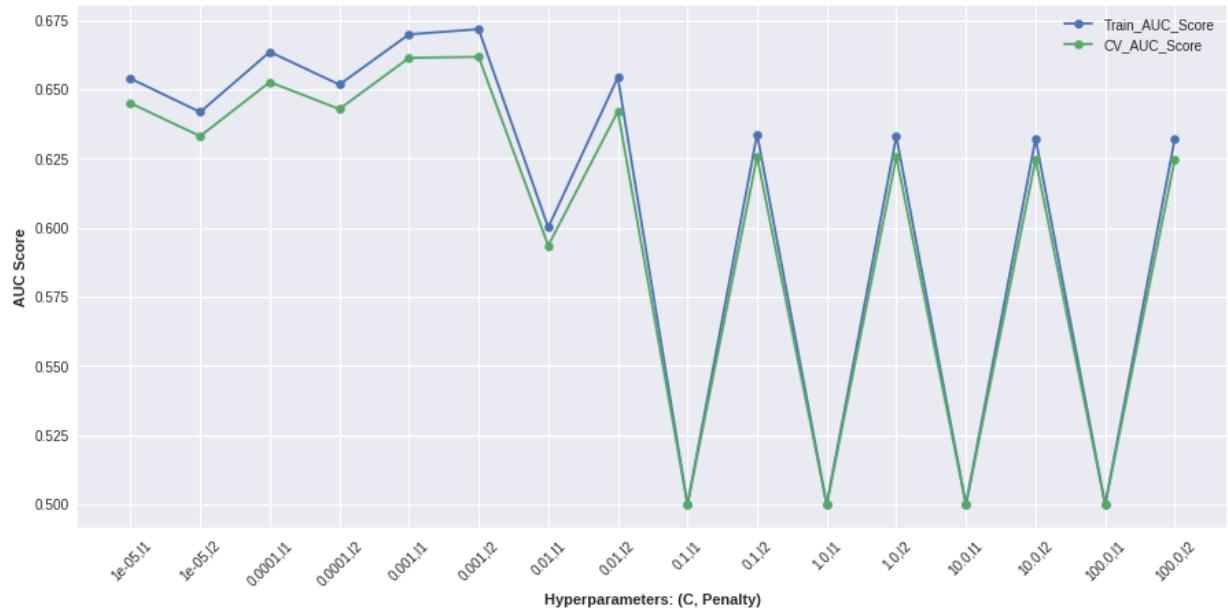
```
In [793]: X_tr.shape,X_cr.shape,X_te.shape
```

```
Out[793]: ((61178, 172), (26220, 172), (21850, 172))
```

```
In [795]: optimal_hyp(X_tr,y_train,X_cr,y_val)
```

```
execution time in minutes: 0.12984533707300822  
execution time in hours: 0
```

```
In [796]: plot_auc(df1)
```



In [797]: `plot_heatmap()`



In [798]: `best_c, p, tr_auc_score, cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]  
print("optimal c:", best_c, "\nRegularizer:", p, "\nCV AUC score:", cv_auc_score)`

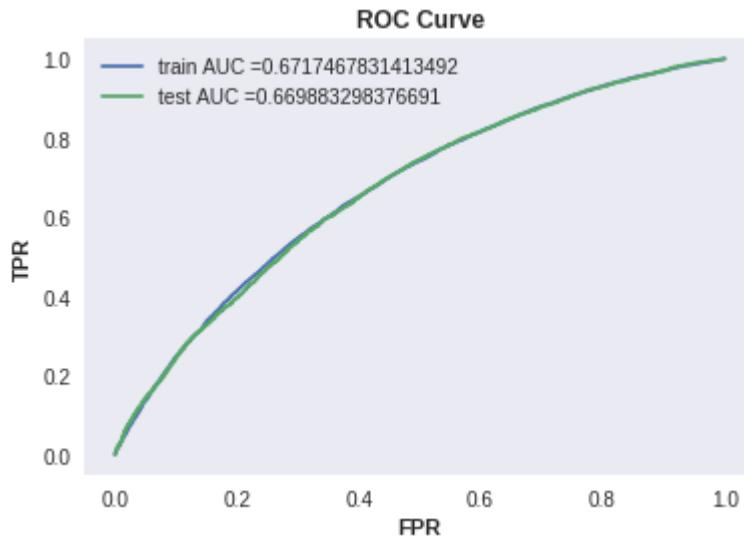
optimal c: 0.001  
Regularizer: 12  
CV AUC score: 0.661743729657827

In [799]: `model = SGDClassifier(class_weight='balanced', alpha=float(best_c), penalty=p, l0  
model.fit(X_tr, y_train)`

Out[799]: `SGDClassifier(alpha=0.001, average=False, class_weight='balanced',  
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,  
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,  
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l2',  
power_t=0.5, random_state=42, shuffle=True, tol=None,  
validation_fraction=0.1, verbose=0, warm_start=False)`

In [0]: `y_pred = model.predict(X_te)`

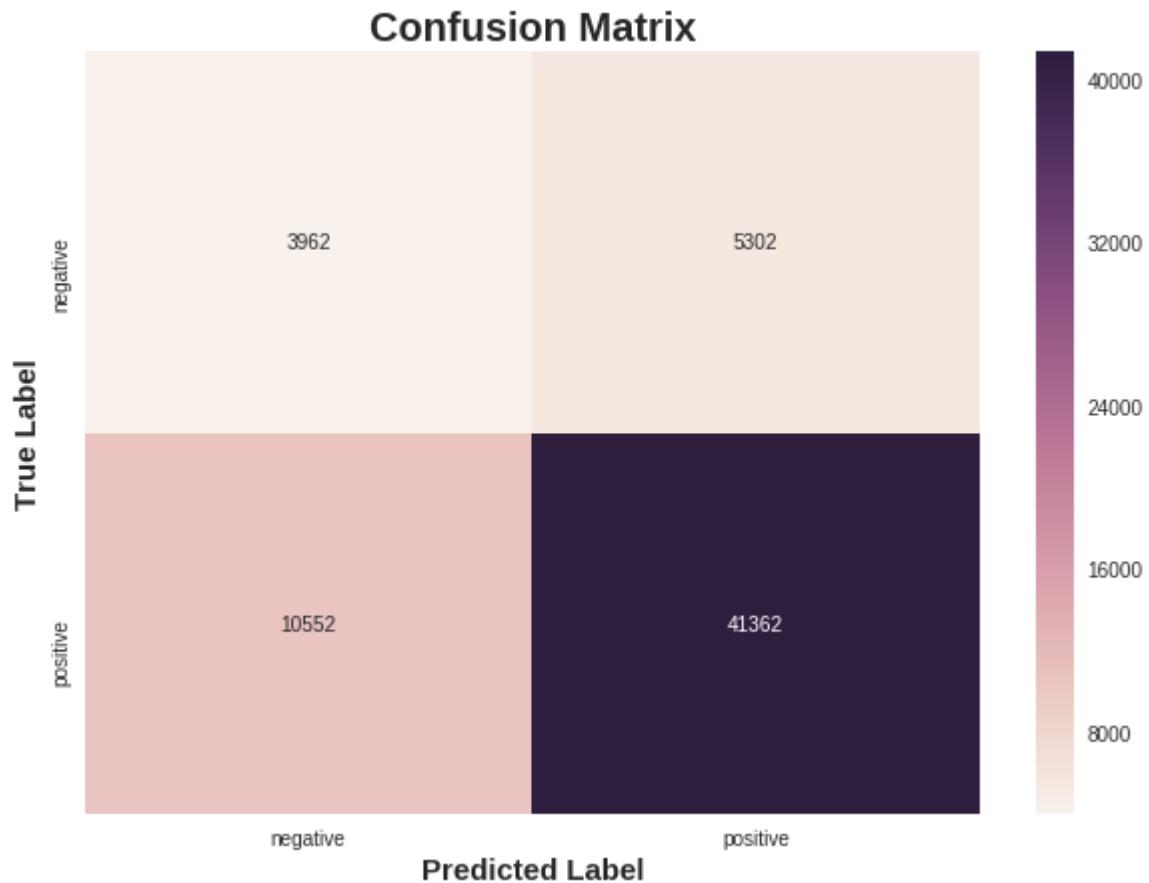
```
In [801]: start = time.time()
error_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
=====
=====
execution time in minutes: 0.010188098748524983
execution time in hours: 0
```

### confusion matrix for train data

```
In [802]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_train), y_train_pred1)
```



True Positives : 41362

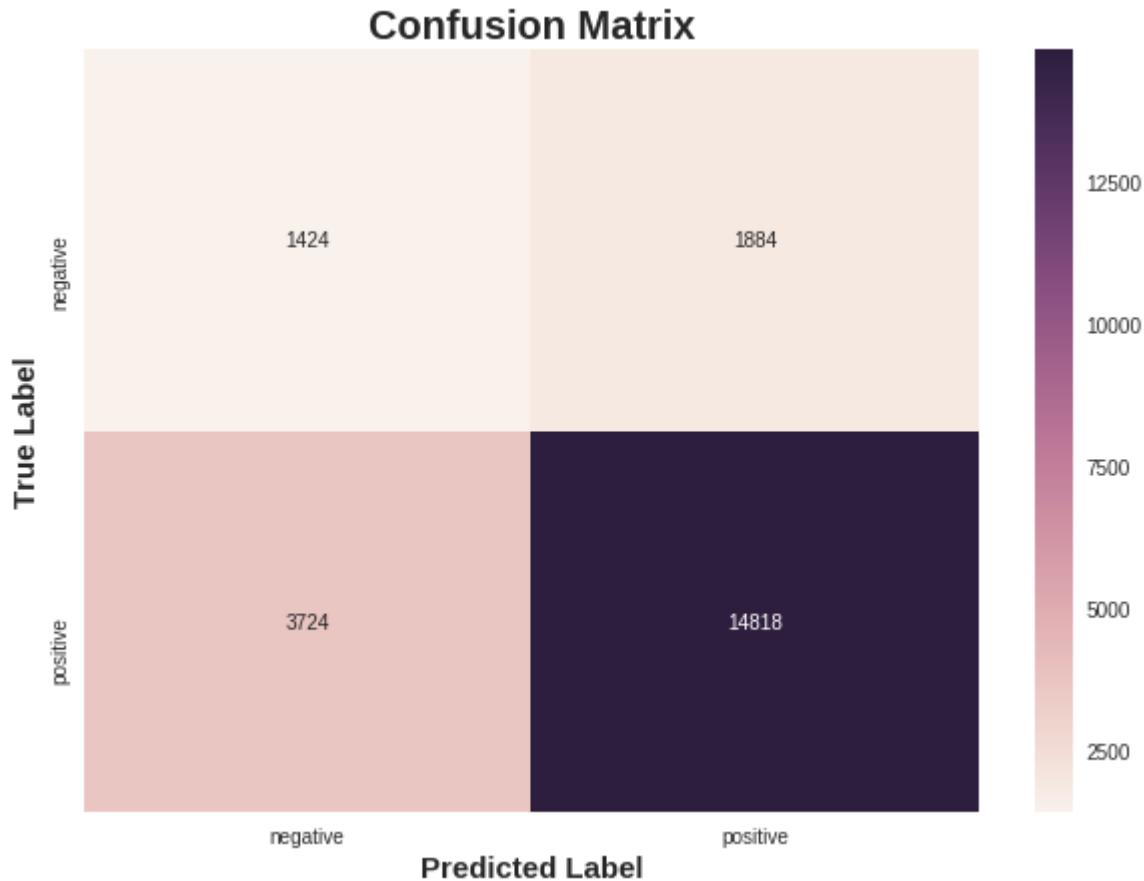
False Positives : 5302

True Negatives : 3962

False Negatives : 10552

**confusion matrix for test data**

```
In [803]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_test), y_pred)
```



True Positives : 14818

False Positives : 1884

True Negatives : 1424

False Negatives : 3724

```
In [805]: # Printing roc_auc score
y_pred = model.decision_function(x_te)
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[805]: 0.669883298376691

## [Task-2] Apply the Support Vector Machines on these features by finding the best hyper parameter

- school\_state : categorical data
- clean\_categories : categorical data
- clean\_subcategories : categorical data
- project\_grade\_category :categorical data
- teacher\_prefix : categorical data
- quantity : numerical data
- teacher\_number\_of\_previously\_posted\_projects : numerical data
- price : numerical data
- sentiment score's of each of the essay : numerical data

- number of words in the title : numerical data
- number of words in the combine essays : numerical data
- Apply TruncatedSVD on TfidfVectorizer of essay text, choose the number of components (n\_components) using elbow method : numerical data

```
In [0]: """
import nltk
from nltk.sentiment.vader import SentimentIntensityAnalyzer
nltk.download('vader_lexicon')
sid = SentimentIntensityAnalyzer()
p = project_data.copy()
sentiments = []
for i in range(p.shape[0]):
    line = p['essay'].iloc[i]
    sentiment = sid.polarity_scores(line)
    sentiments.append([sentiment['neg'], sentiment['pos'], sentiment['neu'], sentiment['compound']])
p[['neg', 'pos', 'neu', 'compound']] = pd.DataFrame(sentiments)
"""

Out[38]: "import nltk\nfrom nltk.sentiment.vader import SentimentIntensityAnalyzer\nnltk.download('vader_lexicon')\nsid = SentimentIntensityAnalyzer()\np = project_data.copy()\nsentiments = []\nfor i in range(p.shape[0]):\n    line = p['essay'].iloc[i]\n    sentiment = sid.polarity_scores(line)\n    sentiments.append([sentiment['neg'], sentiment['pos'], sentiment['neu'], sentiment['compound']])\np[['neg', 'pos', 'neu', 'compound']] = pd.DataFrame(sentiments)\n"
```

```
In [45]: #project_data = p
%cd -
/content/drive/My Drive
```

```
In [0]: project_data= pd.read_pickle('set5')
```

```
In [0]: project_data['words_in_title'] = project_data['project_title'].str.count(' ') + 1
project_data['words_in_essay'] = project_data['essay'].str.count(' ') + 1
project_data['words_in_summary'] = project_data['project_resource_summary'].str.count(' ') + 1
```

```
In [0]: #project_data = project_data[:70000]
```

```
In [0]: project_data_features = project_data.copy()
project_data_features.drop('project_is_approved', axis=1, inplace=True)
y=list(project_data['project_is_approved'])
X_train, X_test, y_train, y_test = train_test_split(project_data_features, y, stratify=y)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, stratify=y)
```

```
In [0]: #X_train.shape, X_val.shape, X_test.shape
```

```
In [0]: neg_standardized_tr,neg_standardized_val, neg_standardized_te = standardize_data()
print()
pos_standardized_tr,pos_standardized_val, pos_standardized_te = standardize_data()
print()
neu_standardized_tr,neu_standardized_val, neu_standardized_te = standardize_data()
print()
comp_standardized_tr,comp_standardized_val, comp_standardized_te = standardize_da
qty_standardized_tr,qty_standardized_val, qty_standardized_te = standardize_data()
price_standardized_tr,price_standardized_val, price_standardized_te = standardize_
print()
project_standardized_tr,project_standardized_val, project_standardized_te = stand
print()
title_standardized_tr, title_standardized_val, title_standardized_te = standardiz
print()
essay_standardized_tr,essay_standardized_val, essay_standardized_te = standardize_
print()
resource_standardized_tr,resource_standardized_val, resource_standardized_te = st
cat_one_hot_tr,cat_one_hot_val,cat_one_hot_te,_ = vectorized_data(X_train,X_val,X_
cat_sub_one_hot_tr,cat_sub_one_hot_val,cat_sub_one_hot_te,_ = vectorized_data(X_t
school_state_dict = create_dict(X_train,'school_state')
teacher_prefix_dict = create_dict(X_train,'teacher_prefix')

state_one_hot_tr,state_one_hot_val,state_one_hot_te,_ = vectorized_data(X_train,X_
teacher_one_hot_tr,teacher_one_hot_val,teacher_one_hot_te,_ = vectorized_data(X_t
grade_dict = create_dict(X_train,'project_grade_category')
grade_one_hot_tr,grade_one_hot_val,grade_one_hot_te,_ = vectorized_data(X_train,X_
one_hot_num_dig_tr = num_hot_encode(X_train,'numerical_digits')
one_hot_num_dig_te = num_hot_encode(X_test,'numerical_digits')
one_hot_num_dig_val = num_hot_encode(X_val,'numerical_digits')
```

```
In [0]: from scipy.sparse import hstack

f_tr = hstack((one_hot_num_dig_tr,grade_one_hot_tr,state_one_hot_tr,cat_one_hot_t
                project_standardized_tr,title_standardized_tr, essay_standardized_
                neg_standardized_tr, pos_standardized_tr, neu_standardized_tr,comp_
f_val = hstack((one_hot_num_dig_val,grade_one_hot_val,state_one_hot_val,cat_one_ho
                project_standardized_val,title_standardized_val, essay_standardized_
                neg_standardized_val, pos_standardized_val, neu_standardized_val,co
f_te = hstack((one_hot_num_dig_te,grade_one_hot_te,state_one_hot_te,cat_one_hot_te
                project_standardized_te,title_standardized_te, essay_standardized_
                neg_standardized_te, pos_standardized_te, neu_standardized_te,comp_
```

```
In [0]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import TruncatedSVD

def tfidf_vec(preprocessed_data_tr, preprocessed_data_val, preprocessed_data_te):
    vectorizer = TfidfVectorizer(min_df=10)
    text_tfidf_tr = vectorizer.fit_transform(preprocessed_data_tr)
    print("Shape of matrix after one hot encoding ", text_tfidf_tr.shape)

    text_tfidf_val = vectorizer.transform(preprocessed_data_val)
    print("Shape of matrix after one hot encoding ", text_tfidf_val.shape)

    text_tfidf_te = vectorizer.transform(preprocessed_data_te)
    print("Shape of matrix after one hot encoding ", text_tfidf_te.shape)
    return text_tfidf_tr, text_tfidf_val, text_tfidf_te
```

```
In [54]: tfidf_vec_essay_tr, tfidf_vec_essay_val, tfidf_vec_essay_te = tfidf_vec(textpreproc...
```

100% | ██████████ | 61178/61178 [00:37<00:00, 1610.82it/s]  
100% | ██████████ | 26220/26220 [00:16<00:00, 1594.71it/s]  
100% | ██████████ | 21850/21850 [00:13<00:00, 1599.75it/s]

Shape of matrix after one hot encoding (61178, 13194)  
Shape of matrix after one hot encoding (26220, 13194)

6% | | 3479/61178 [00:00<00:01, 34783.48it/s]

Shape of matrix after one hot encoding (21850, 13194)

100% | ██████████ | 61178/61178 [00:01<00:00, 35226.80it/s]  
100% | ██████████ | 26220/26220 [00:16<00:00, 1608.36it/s]  
100% | ██████████ | 21850/21850 [00:00<00:00, 34455.40it/s]

Shape of matrix after one hot encoding (61178, 2265)

0% | | 0/61178 [00:00<?, ?it/s]

Shape of matrix after one hot encoding (26220, 2265)  
Shape of matrix after one hot encoding (21850, 2265)

100% | ██████████ | 61178/61178 [00:03<00:00, 15334.73it/s]  
100% | ██████████ | 26220/26220 [00:16<00:00, 1607.81it/s]  
100% | ██████████ | 21850/21850 [00:01<00:00, 15262.35it/s]

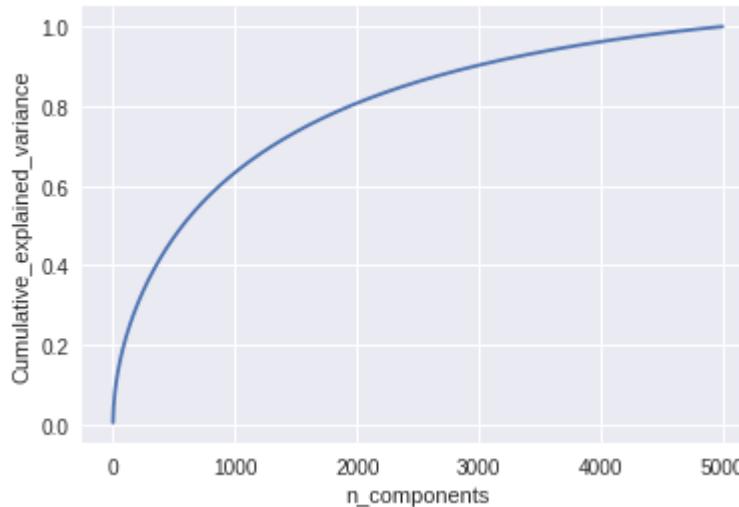
Shape of matrix after one hot encoding (61178, 4351)  
Shape of matrix after one hot encoding (26220, 4351)  
Shape of matrix after one hot encoding (21850, 4351)

```
In [0]: %time
# https://chrisalbon.com/machine_learning/feature_engineering/select_best_number_of_
X = tfidf_vec_essay_tr
# Create and run an TSVD with one less than number of features
tsvd = TruncatedSVD(n_components=4999)
X_tsvd = tsvd.fit(X) #####
```

```
In [0]: # List of explained variances
tsvd_var_ratios = tsvd.explained_variance_ratio_
```

```
In [85]: percentage_var_explained = tsvd.explained_variance_ / np.sum(tsvd.explained_variance)
cum_var_explained = np.cumsum(percentage_var_explained)

plt.clf()
plt.plot(cum_var_explained)
plt.xlabel('n_components')
plt.ylabel('Cumulative_explained_variance')
plt.show()
```



```
In [0]: # Create a function
def select_n_components(var_ratio, goal_var: float) -> int:
    # Set initial variance explained so far
    total_variance = 0.0

    # Set initial number of features
    n_components = 0

    # For the explained variance of each feature:
    for explained_variance in var_ratio:

        # Add the explained variance to the total
        total_variance += explained_variance

        # Add one to the number of components
        n_components += 1

        # If we reach our goal level of explained variance
        if total_variance >= goal_var:
            # End the loop
            break

    # Return the number of components
    return n_components
```

```
In [0]: # Run function
n=select_n_components(tsvd_var_ratios, 0.85)
```

```
In [0]: # Run function
ncomp=select_n_components(tsvd_var_ratios, 0.95)
```

```
In [0]: X = tfidf_vec_essay_tr
# Create and run an TSVD with one less than number of features
tsvd = TruncatedSVD(n_components=n)
X_tsvd = tsvd.fit(X)
```

```
In [0]: tfidf_essay_tr = tsvd.transform(X)
tfidf_essay_val = tsvd.transform(tfidf_vec_essay_val)
tfidf_essay_te = tsvd.transform(tfidf_vec_essay_te)
```

```
In [0]: from scipy.sparse import hstack

def hstack_data(f1_tr, f1_cr, f1_te):
    X_tr = hstack((f1_tr, f1_tr)).tocsr()
    X_cr = hstack((f1_cr, f1_cr)).tocsr()
    X_te = hstack((f1_te, f1_te)).tocsr()
    return X_tr,X_cr,X_te
```

```
In [0]: X_tr, X_cr, X_te = hstack_data(tfidf_essay_tr, tfidf_essay_val, tfidf_essay_te)
```

```
In [59]: xtrain = pd.DataFrame(X_tr.toarray())
xtrain['label'] = y_train
final = xtrain
final.sort_values(by='label', ascending=False, inplace=True)
n_zeros = final[final['label']==0]
n_ones = final[final['label']==1]
diff = n_ones.shape[0] - n_zeros.shape[1]
n_zeros.shape, n_ones.shape
```

Out[59]: ((9264, 691), (51914, 691))

```
In [0]: n_ones = n_ones[:40000]
from sklearn.utils import resample
final_neg = n_zeros
final_neg_upsampled = resample(final_neg, replace=True, n_samples=diff, random_st
final = pd.concat([n_ones, final_neg_upsampled])
```

```
In [61]: final.shape
```

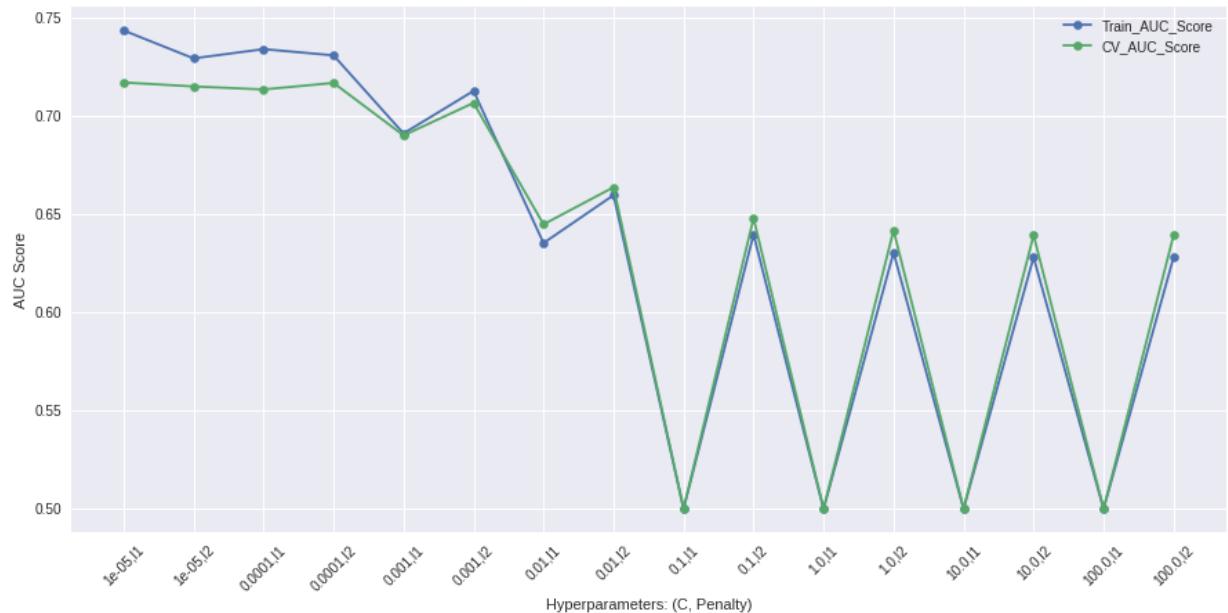
Out[61]: (91223, 691)

```
In [0]: import scipy
y_train = final['label']
X_tr = final.drop('label', axis=1)
X_tr = scipy.sparse.csr_matrix(X_tr.values)
```

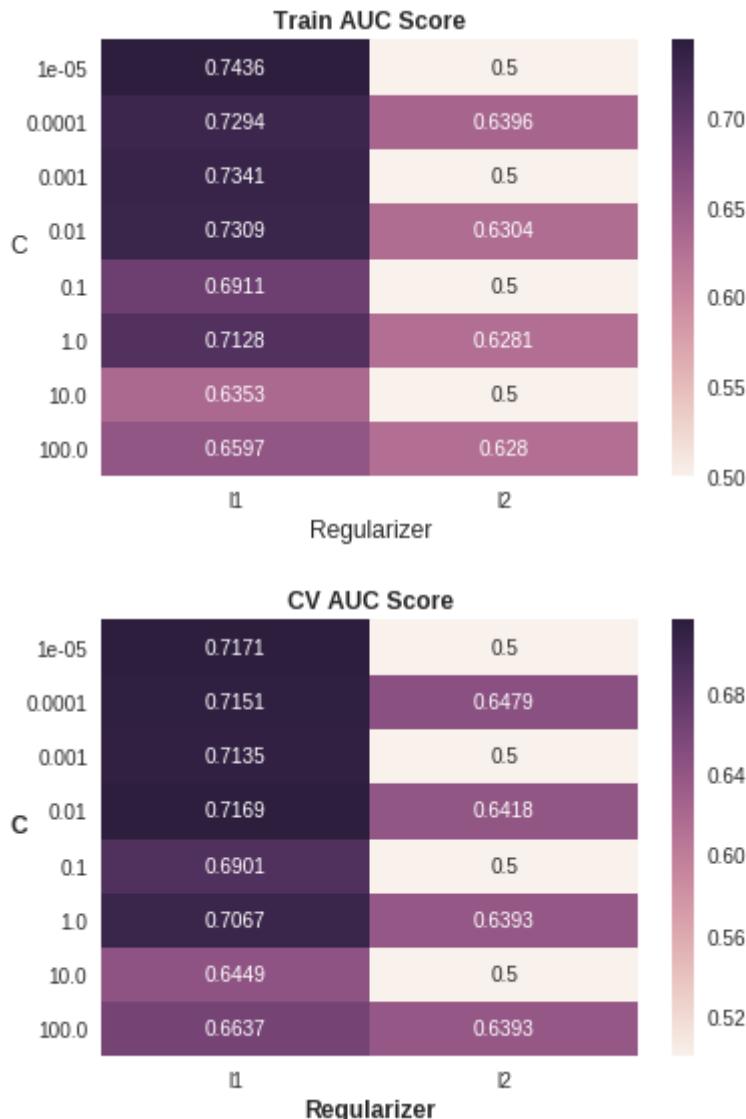
```
In [63]: optimal_hyp(X_tr,y_train,X_cr,y_val)
```

execution time in minutes: 0.8826952576637268  
 execution time in hours: 0

```
In [64]: plot_auc(df1)
```



In [65]: `plot_heatmap()`



In [66]: `best_c, p, tr_auc_score, cv_auc_score = df1.iloc[df1.CV_AUC_Score.argmax()]  
print("optimal c:", best_c, "\nRegularizer:", p, "\nCV AUC score:", cv_auc_score)`

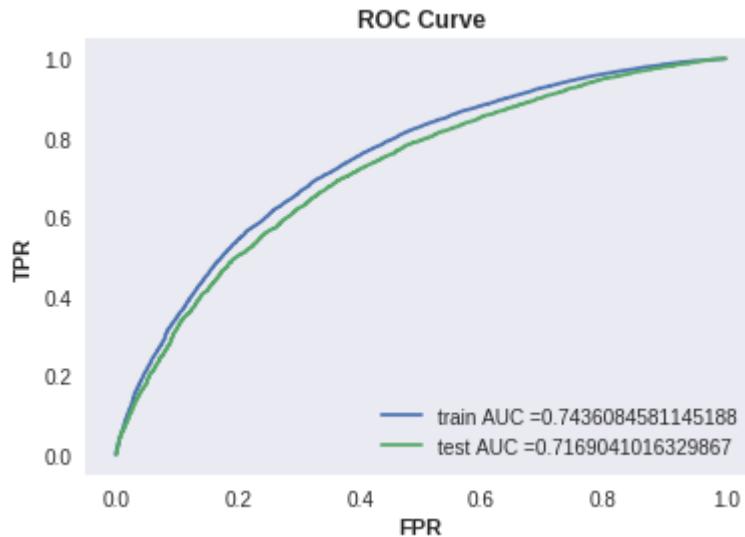
optimal c: 1e-05  
Regularizer: l1  
CV AUC score: 0.7171011122746441

In [67]: `model = SGDClassifier(class_weight='balanced', alpha=float(best_c), penalty=p, l0  
model.fit(X_tr, y_train)`

Out[67]: `SGDClassifier(alpha=1e-05, average=False, class_weight='balanced',  
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,  
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,  
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l1',  
power_t=0.5, random_state=42, shuffle=True, tol=None,  
validation_fraction=0.1, verbose=0, warm_start=False)`

In [0]: `y_pred = model.predict(X_te)`

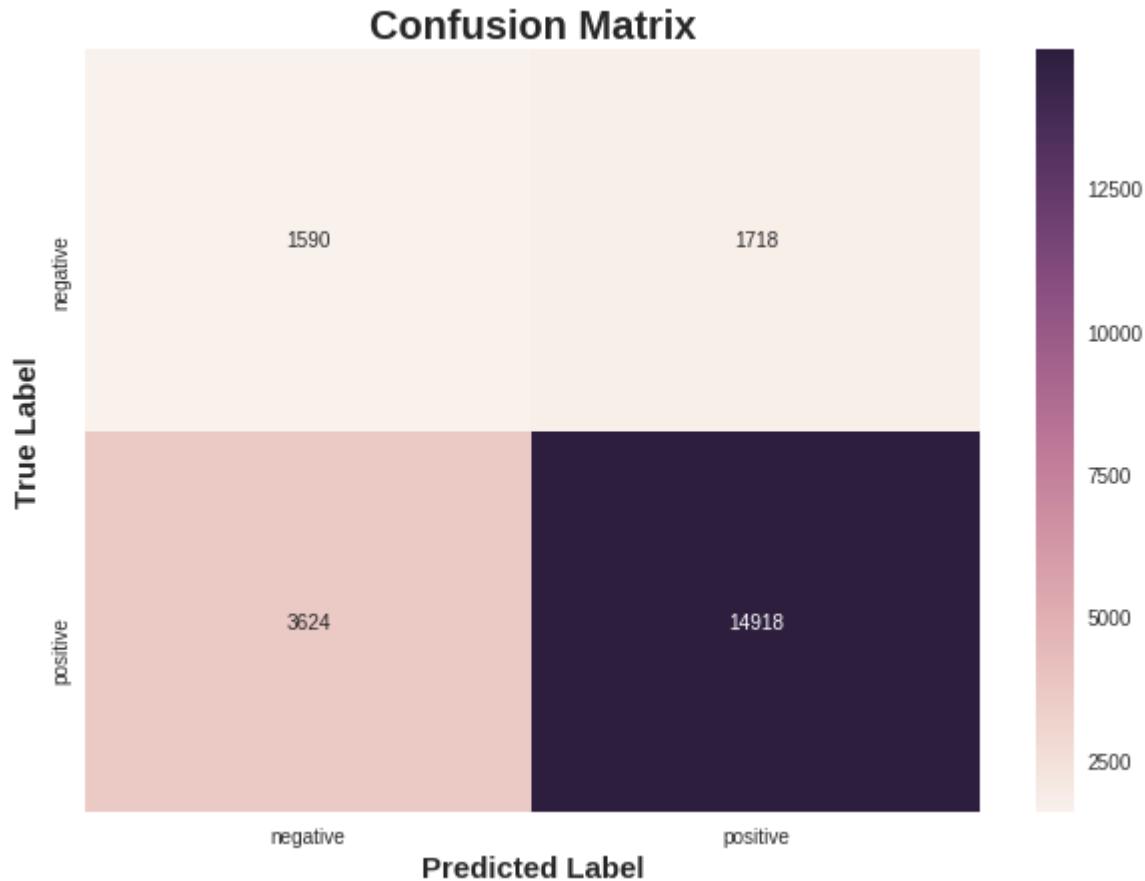
```
In [69]: start = time.time()
error_plot(X_tr,X_te,y_train,y_test)
end = time.time()
minutes = float((end - start)/60)
print("execution time in minutes:",minutes)
print("execution time in hours:",int(minutes/60))
```



```
=====
=====
execution time in minutes: 0.05978336334228516
execution time in hours: 0
```

### confusion matrix for test data

```
In [72]: labels = ["negative","positive"]
get_confusion_matrix_values(np.array(y_test), y_pred)
```



```
In [73]: # Printing roc auc score
y_pred = model.decision_function(X_te)
roc_auc_score(y_true=y_test, y_score=y_pred)
```

Out[73]: 0.7169041016329867

## Summary

```
In [86]: from prettytable import PrettyTable
import sys
sys.stdout.write("\033[1;30m")

x = PrettyTable()
x.field_names = ["Vectorizer", "C", "Penalty", "AUC"]

x.add_row(["BOW", 0.01, 'l2', 0.7289])
x.add_row(["TFIDF", 0.001, 'l2', 0.7234])
x.add_row(["AVGW2V", 0.0001, 'l1', 0.6654])
x.add_row(["TFIDFW2V", 0.001, 'l2', 0.6698])
x.add_row(["Task2", 0.00001, 'l1', 0.7169])

print(x)
```

Vectorizer	C	Penalty	AUC
BOW	0.01	l2	0.7289
TFIDF	0.001	l2	0.7234
AVGW2V	0.0001	l1	0.6654
TFIDFW2V	0.001	l2	0.6698
Task2	1e-05	l1	0.7169